

Modern C++ Programming

4. BASIC CONCEPTS III

Federico Busato

University of Verona, Dept. of Computer Science
2020, v3.01



1 Memory Management: Heap and Stack

- `new`, `delete`
- Memory Leak
- Stack Allocation

2 Initialization

3 Pointers and References

- Pointer
- Address-of operator `&`
- Reference

4 sizeof Operator

5 const and constexpr

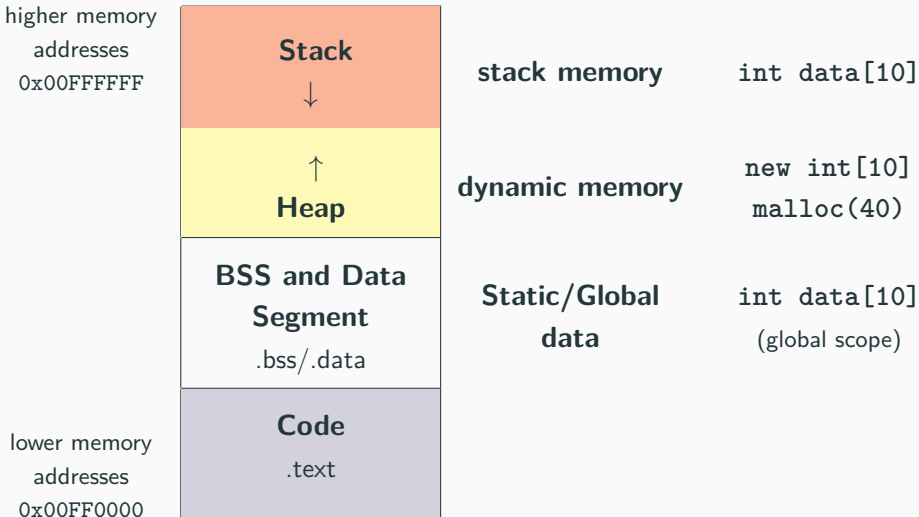
- const
- constexpr

6 Explicit Type Conversion

- static_cast, const_cast, reinterpret_cast
- Type Punning
- Brace Initialization

Memory Management: Heap and Stack

Process Address Space



new, delete

`new` / `new[]` and `delete` / `delete[]` are C++ *keywords* that perform dynamic memory allocation/deallocation, and object construction/ destruction at runtime

`malloc` and `free` are C functions and they allocate and free *memory blocks* (expressed in bytes)

Example:

```
int* array = new int[10]; // C: (int*) malloc(10 * sizeof(int))
delete[] array;           // C: free(array)
```

new, delete Advantages

- **Language keywords**, not functions → *safer*
- **Return type**: `new` returns exact data type, while `malloc()` returns `void*`
- **Failure**: `new` throws an *exception*, while `malloc()` returns a `NULL` pointer → *it cannot be ignored*
- **Allocated bytes**: The size of the allocated memory is calculated by the compiler for `new`, while the user must take care of manually calculate the size for `malloc()`
- **Initialization**: `new` can be used to initialize an object or a set of objects

Dynamic Allocation

- Allocate a single element

```
int* value = (int*) malloc(sizeof(int)); // C
int* value = new int;                  // C++
```

- Allocate N elements

```
int* array = (int*) malloc(N * sizeof(int)); // C
int* array = new int[N];                    // C++
```

- Allocate and zero-initialize N elements

```
int* array = (int*) calloc(N * sizeof(int)); // C
int* array = new int[N]();                  // C++
```

- Allocate N structures

```
MyStruct* array = (int*) malloc(N * sizeof(MyStruct)); // C
MyStruct* array = new MyStruct[N];                    // C++
```


Dynamic Deallocation

- Deallocate a single element

```
int* value = (int*) malloc(sizeof(int)); // C  
free(value);
```

```
int* value = new int; // C++  
delete value;
```

- Deallocate N elements

```
int* value = (int*) malloc(N * sizeof(int)); // C  
free(value);
```

```
int* value = new int[N]; // C++  
delete[] value;
```

Fundamental rules:

- Each object allocated with `new` must be deallocated with `delete`
- Each object allocated with `new[]` must be deallocated with `delete[]`

Mixing `new`, `new[]`, `malloc` with something different from their counterparts leads to *undefined behavior*

`delete` and `delete[]` applied to `NULL`/`nullptr` pointers do not produce errors (same as `free`)

Memory Leak

Memory Leak

A **memory leak** is a dynamically allocated entity in heap memory that is no longer used by the program, but still maintained overall its execution

Problems:

- Illegal memory accesses → segmentation fault
- Undefined values → segmentation fault
- Additional memory consumption

```
int main() {  
    int* array = new int[10];  
    array      = nullptr; // memory leak!!  
} // the memory can no longer be deallocated!!
```

Note: the memory leaks are especially difficult to detect in complex code and when objects are widely used

Wild and Dangling Pointers

Wild pointer:

```
int main() {  
    int* ptr;    // wild pointer: Where will this pointer points?  
    ...         // solution: always initialize a pointer  
}
```

Dangling pointer:

```
int main() {  
    int* array = new int[10];  
    delete[] array; // ok -> "array" now is a dangling pointer  
    delete[] array; // double free or corruption!!  
    // program aborted, the value of "array" is not null  
}
```

Solution:

```
int main() {  
    int* array = new int[10];  
    delete[] array; // ok -> "array" now is a dangling pointer  
    array = nullptr; // no more dangling pointer  
    delete[] array; // ok, no side effect  
}
```

Stack Allocation

A local variable is either in stack memory or CPU registers

Every object which resides in the stack is not valid outside the current scope!!

```
int* wrongFunction() {  
    int A[3] = {1, 2, 3};  
    return A;  
}  
  
int main() {  
    int* ptr = wrongFunction();  
    cout << ptr[0]; // Illegal memory access!!  
}
```

The organization of stack memory enables much higher performance. On the other hand, this memory space is limited!!

It is $\approx 8MB$ on linux by default

2D Memory Allocation

Easy on stack:

```
int A[3][4];
```

Dynamic Memory 2D allocation/free:

```
int** A = new int*[3];  
for (int i = 0; i < 3; i++)  
    A[i] = new int[4];  
  
for (int i = 0; i < 3; i++)  
    delete[] A[i];  
delete[] A;
```

Dynamic memory 2D allocation/free C++11:

```
auto A = new int[3][4];    // allocate 3 objects of size int[4]  
int n = 3;                // dynamic value  
auto B = new int[n][4];    // ok  
// auto C = new int[n][n]; // compile error!!  
delete[] A;               // same for B, C
```

Data and BSS Segment

```
int data[]          = {1, 2}; // DATA segment memory
int big_data[1000000] = {};    // BSS segment memory
                                // (zero-initialized)

int main() {
    int A[] = {1, 2, 3}; // stack memory
}
```

Data/BSS (Block Started by Symbol) segments are larger than stack memory (max \approx 1GB in general) but slower

Initialization

Variable Initialization

C++03:

```
int a;           // undefined
int a1 = 2;      // copy initialization
int a2(2);       // direct initialization
int a3 = {2};    // copy list initialization
int a4 = int(2); // copy initialization
int a5 = int();  // copy initialization (zero-initialization)
// int a6();     // a6 is a function
```

C++11 (Uniform Initialization):

```
int b1{2};       // direct list initialization
int b2 = {2};    // copy list initialization

// Zero-initialization / Default-Initialization
int b3 = {};     // copy list initialization, default value
int b4{};        // direct list initialization, default value
```

Stack Array Initialization

One dimension:

```
int A[3] = {1, 2, 3}; // explicit size
int B[] = {1, 2, 3}; // implicit size
char C[] = "abcd";    // implicit size
int C[3] = {1, 2};    // C[2] = 0 -> zero/default value

int D[4] = {0}; // all values of D are initialized to 0
int E[3] = {};  // all values of E are initialized to 0 (C++11)
int F[3] {};    // all values of F are initialized to 0 (C++11)
```

Two dimensions:

```
int G[][2] = { {1,2}, {3,4}, {5,6} }; // ok
int H[][2] = { 1, 2, 3, 4 };           // ok
// the type of G and H is an array of type int[]
// int F[][] = ...;                     // compile error!!
// int G[2][] = ...;                    // compile error!!
```

Rules for Zero/Default Initialization:

- An object with **dynamic** storage duration (heap) has indeterminate value
- An object whose initializer is an **empty set of parentheses** `{ }` is zero or default initialized (**C++11**)

Dynamic Initialization

C++03:

```
int* a1 = new int;           // undefined
int* a2 = new int();         // zero-initialization
int* a3 = new int(4);        // allocate a single value equal to 4
int* a4 = new int[3];        // allocate 4 elements with undefined values
int* a5 = new int[4]();      // allocate 4 elements zero-initialized
```

C++11:

```
int* b1 = new int[4]{};      // allocate 4 elements zero-initialized
int* b2 = new int[4]{1, 2};  // set first, second, zero-initialized
```

Pointers and References

Pointers and Pointer Dereferencing

Pointer

A **pointer** is a value referring to a location in memory

Pointer Dereferencing

Pointer **dereferencing** means obtaining the value stored in at the location referred to the pointer

```
int* ptr1 = new int;  
*ptr1     = 4;      // dereferencing (assignment)  
int a     = *ptr1;  // dereferencing (get value)
```

Common error:

```
int *ptr1, ptr2; // one pointer and one integer!!  
int *ptr1, *ptr2; // ok, two pointers
```

void Pointer (Generic Pointer)

Instead of declaring different types of pointer variable it is possible to declare single pointer variable which can act as any pointer types

- A `void*` can be assigned to another `void*`
- `void*` can be compared for equality and inequality
- A `void*` can be explicitly converted to another type
- Other operations would be unsafe because the compiler cannot know what kind of object is really pointed to. Consequently, other operations result in compile-time errors

```
cout << (sizeof(void*) == sizeof(int*)); // print true

int array[] = { 2, 3, 4 };
void* ptr = array;
cout << *array;           // print 2
// cout << *ptr;          // compile error!!
cout << *((int*) ptr);    // print 2
// void* ptr2 = ptr + 2;  // compile error!!
```

Address-of operator &

The **address-of operator** (&) returns the address of a variable

```
int a = 3;
int* b = &a; // address-of operator,
             // 'b' is equal to the address of 'a'
a++;
cout << *b; // print 4;
```

To not confuse with **Reference syntax**: `T& var = ...`

```
int array[4];
// &array is a pointer to an array of size 4
int size1 = (&array)[1] - array;
int size2 = *(&array + 1) - array;
cout << size1; // print 4
cout << size2; // print 4
```


1 + 1 \neq 2 : Pointer Arithmetic

Pointer syntax:

`ptr[i]` is equal to `*(ptr + i)`

Pointer arithmetic rule:

`address(ptr + i) = address(ptr) + (sizeof(T) * i)`

where T is the type of elements pointed by ptr

Example:

```
int array[4] = {1, 2, 3, 4};  
cout << array[1];           // print 2  
cout << *(array + 1);       // print 2  
cout << array;               // print 0xFFFFFFF2  
cout << array + 1;          // print 0xFFFFFFF6!!
```

```
char arr[3] = "abc"
```

value	address	
'a'	0x0	\leftarrow arr[0]
'b'	0x1	\leftarrow arr[1]
'c'	0x2	\leftarrow arr[2]

```
int arr[3] = {4,5,6}
```

value	address	
4	0x0	\leftarrow arr[0]
	0x1	
	0x2	
	0x3	
5	4	\leftarrow arr[1]
	0x5	
	0x6	
	0x7	
	0x8	\leftarrow arr[2]

Reference

A variable **reference** is an **alias**, namely another name for an already existing variable. Both variable and variable reference can be applied to refer the value of the variable

- A pointer has its own memory address and size on the stack, reference shares the **same memory address** (with the original variable)
- References are internally implemented as *pointer*, but the compiler treats them in a very different way

References are safer than pointers:

- References cannot have NULL value. You must always be able to assume that a reference is connected to a legitimate storage
- References cannot be changed. Once a reference is initialized to an object, it cannot be changed to refer to another object (Pointers can be pointed to another object at any time)
- References must be initialized when they are created (Pointers can be initialized at any time)

Reference (Examples)

Reference syntax: `T& var = ...`

```
//int& d;    // reference. compile error!! no initialization  
int c = 2;  
int& e = c; // reference. ok valid initialization  
e++;      // increment  
cout << c; // print 3
```

```
int a = 3;  
int* b = &a; // pointer  
int* c = &a; // pointer  
b++;      // change the value of the pointer 'b'  
*c++;     // change the value of 'a'  
  
int& c = a; // reference  
c++;      // change the value of 'a'
```

Reference (Function Arguments)

Reference vs. pointer arguments:

```
void f(int* value) {} // value may be a nullptr
void g(int& value) {} // value is never a nullptr

int a = 3;
f(&a); // ok
g(a); // ok
//g(3); // compile error!! "3" is not a reference of something
```

References can be use to indicate fixed size arrays:

```
f(int (&array)[3]) {} // accepts only arrays of size 3
                      // f(int array[]) accepts any size

int A[3], B[4];
int* C = A;
//-----
f(A); // ok
// f(B); // compile error!! B has size 4
// f(C); // compile error!! C is a pointer
```

Reference (Arrays)

```
int A[4];  
int (&B)[4] = A;      // ok, reference to array  
int C[10][3];  
int (&D)[10][3] = C; // ok, reference to 2D array  
  
auto c = new int[3][4]; // type is int (*)[4]  
// read as "pointer to arrays of 4 int"  
// int (&d)[3][4] = c; // compile error!!  
// int (*e)[3]     = c; // compile error!!  
int (*f)[4] = c;      // ok
```

Reference:

[1] www3.ntu.edu.sg/home/ehchua/programming/cpp/cp4_PointerReference.html

Reference and struct

- The **dot** (.) operator is applied to local objects and references
- The **arrow** operator (->) is used with a pointer to an object

```
#include <iostream>
struct A {
    int x = 3;
};

int main() {
    A obj;

    A* p = &obj;    // pointer
    p->x;            // arrow syntax

    A& ref = obj;    // reference
    cout << obj.x;    // dot syntax
    cout << ref.x;    // dot syntax
}
```

sizeof Operator

sizeof operator

sizeof

The `sizeof` is a compile-time operator that determines the size, in bytes, of a variable or data type

- `sizeof` returns a value of type `size_t`
- `sizeof(incomplete type)` produces compile error, e.g. `void`
- `sizeof(bitfield)` produces compile error
- `sizeof(anything)` never returns 0, except for array of size 0
- `sizeof(char)` always returns 1
- When applied to structures, it also takes into account padding
- When applied to a reference, the result is the size of the referenced type

```
sizeof(int);    // 4
sizeof(int*)    // 8 on a 64-bit OS
sizeof(void*)   // 8 on a 64-bit OS
sizeof(size_t)  // 8 on a 64-bit OS
```

```
int f(int[] array) {           // dangerous!!
    cout << sizeof(array);
}

int array1[10];
int* array2 = new int[10];
cout << sizeof(array1); // print sizeof(int) * 10 = 40
cout << sizeof(array2); // print sizeof(int*) = 8 (64-bit)
f(array1);                // print 8
```

```
struct B {  
    int x;  
    char y;  
};  
  
struct C : B { // C extends B  
    short z;  
};  
  
//-----  
sizeof(B);      // 8 : 4 + 1 (+ 3) (padding)  
sizeof(C);      // 12 : sizeof(B) + 2 (+ 2) (padding)  
  
struct A {};  
sizeof(A);      // 1 : sizeof never return 0 (except for arrays)
```

```
char a;  
char& b = a;  
sizeof(&a);    // 8 in a 64-bit OS (pointer)  
sizeof(b);     // 1, i.e. sizeof(char)  
              // NOTE: a reference is not a pointer  
  
struct A {};  
A array1[10];  
sizeof(array1); // 1 : array of empty structures  
  
int array2[0];  
sizeof(array2); // 0 : special case
```

const and constexpr

const Keyword

const keyword

The `const` keyword indicates objects never changing value after their initialization (they must be initialized when declared)

`const` variables are evaluated at compile-time value if the right expression is also evaluated at compile-time

```
int size = 3;
int A[size] = {1, 2, 3}; // Technically possible (size is dynamic)
                        // But NOT approved by the C++ standard

const int SIZE = 3;
// SIZE = 4;           // compile error!! (SIZE is const)
int B[SIZE] = {1, 2, 3}; // ok

const int size2 = size;
int B[size2] = {1, 2, 3}; // BAD programming!! size is not const
// (some compilers allow variable size stack array -> dangerous!!)
```

Constness rules:

- `int* → const int*`
- `const int* ↯ int*`

```
int f1(const int* array) { // the values of the array cannot be  
    ... // modified  
}
```

```
int f2(int* array) {}
```

```
int* ptr = new int[3];  
const int* c_ptr = new int[3];  
f1(ptr); // ok  
f2(ptr); // ok  
f1(c_ptr); // ok  
// f2(c_ptr); // compile error!!
```

```
void g(const int) { // pass-by-value combined with 'const'  
    ... // note: it is not useful because the value  
} // is copied
```

- `int*` pointer to `int`
 - The value of the pointer can be modified
 - The elements referred by the pointer can be modified
- `const int*` pointer to `const int`. Read as `(const int)*`
 - The value of the pointer can be modified
 - The elements referred by the pointer cannot be modified
- `int *const` const pointer to `int`
 - The value of the pointer cannot be modified
 - The elements referred by the pointer can be modified
- `const int *const` const pointer to `const int`
 - The value of the pointer cannot be modified
 - The elements referred by the pointer cannot be modified

Note: `const int*` is equal to `int const*`

Tip: pointer types should be read from right to left

Common error: adding `const` to a pointer is not the same as adding `const` to a type alias of a pointer

```
using ptr_t      = int*;
using const_ptr_t = const int*;

void f1(const int* ptr) {
    // ptr[0] = 0;          // not allowed: pointer to const objects
    ptr      = nullptr; // allowed
}

void f3(const_ptr_t ptr) { // same as before
    // ptr[0] = 0;          // not allowed: pointer to const objects
    ptr      = nullptr; // allowed
}

void f2(const ptr_t ptr) { // warning!!
    ptr[0] = 0;          // allowed
    // ptr    = nullptr; // not allowed: const pointer to
    //          // modifiable objects
}
```

constexpr (function)

C++11/C++14/C++17 guarantees compile-time evaluation of an function as long as all its arguments are constant

- C++11: `constexpr` must contain exactly one `return` statement and it must not contain loops or switch
- C++14: `constexpr` has no restrictions

constexpr (variable)

C++11/C++14/C++17 `constexpr` variables are evaluated at compile-time

- `const` guarantees the value of a variable to be fixed overall the execution of the program
- `constexpr` tells the compiler that the expression results is at compile-time. `constexpr` value implies `const`

```
const int v1 = 3;      // compile-time evaluation
const int v2 = v1 * 2; // compile-time evaluation

int      a  = 3;      // "a" is dynamic
const int v3 = a;      // run-time evaluation!!

constexpr c1 = v1;     // ok
// constexpr c2 = v3; // compile error!!
```

```
constexpr int square(int value) {
    return value * value;
}

square(4); // compile-time evaluation

int a = 4; // "a" is dynamic
square(a); // run-time evaluation
```

if constexpr

C++17 introduces `if constexpr` feature which allows *conditionally* compiling code based on a *compile-time* value

It is an `if` statement where the branch is chosen at compile-time (similarly to the `#if` preprocessor)

```
void f() {  
    if constexpr (true)  
        std::cout << "compile!";  
    else  
        std::cout << "error!"; // never compiled  
}
```

constexpr example

```
constexpr int fib(int n) {  
    return (n == 0 || n == 1) ? 1 : fib(n - 1) + fib(n - 2);  
}  
  
int main() {  
    if constexpr (sizeof(void*) == 8)  
        return fib(5);  
    else  
        return fib(3);  
}
```

Generated assembly code (x64 OS):

```
main:  
    mov eax, 8  
    ret
```

Explicit Type Conversion

Old style cast (type) value

New style cast:

- `static_cast` does compile-time (not run-time) checking of the types involved In many situations, this can make it the safest type of cast, as it provides the least room for accidental/unsafe conversions between various types

- `reinterpret_cast`

`reinterpret_cast<T*>(v)` equal to `(T*) v`

`reinterpret_cast<T&>(v)` equal to `*((T*) &v)`

- `const_cast` may be used to cast away (remove) constness or volatility

Static cast vs. old style cast:

```
char a[] = {1, 2, 3, 4};  
int* b = (int*) a;           // ok  
cout << b[0];                // print 67305985 not 1!!  
// int* c = static_cast<int*>(a); // compile error!! unsafe conversion
```

Const cast:

```
const int a = 5;  
const_cast<int>(a) = 3; // ok
```

Reinterpret cast: (bit-level conversion)

```
float b = 3.0f;  
// bit representation of b: 01000000010000000000000000000000  
int c = reinterpret_cast<int&>(b);  
// bit representation of c: 01000000010000000000000000000000  
int a[3][4]; // array reshaping example  
int (&b)[2][6] = reinterpret_cast<int (&)[2][6]>(a);  
int (*c)[6] = reinterpret_cast<int (*)[6]>(a);
```


Type Punning

Pointer Aliasing

One pointer **aliases** another when they both point to the same memory location

Type Punning

Type punning refers to circumvent the type system of a programming language to achieve an effect that would be difficult or impossible to achieve within the bounds of the formal language

```
bool is_negativeA(float x) {  
    return x < 0.0;  
}  
  
bool is_negativeB(float x) {  
    unsigned int* ui = (unsigned int *) &x; // gcc warning:  
    return (*ui) & 0x80000000;              // -Wstrict-aliasing  
} // this could lead to undefined behavior
```

Brace Initialization

C++11 The **brace initialization** can be also used to *safely* convert arithmetic types, preventing implicit *narrowing*, i.e potential value loss

The syntax is also more concise than modern casts

```
int      x1 = -1;
int64_t  x2 = int64_t{-1}; // ok
//int64_t x3 = uint64_t{-1}; // compile error!! unsafe

int64_t  y1 = int64_t{x1}; // ok (only GCC, not clang)
//uint64_t y2 = uint64_t{x1}; // compile error!! unsafe

uint64_t z1 = static_cast<uint64_t>(-1); // ok
uint64_t z2 = static_cast<uint64_t>(x1); // ok

float    f1 = float{10e30}; // ok
//float f2 = float{10e40}; // compile error!! unsafe
```