Modern C++ Programming

4. Basic Concepts III

Federico Busato

University of Verona, Dept. of Computer Science
2020, v3.03

Table of Context

Table of Context

Memory Management: Heap and Stack

Process Address Space

higher memory
addresses
0x00FFFFFF

Stack ↓	**stack memory** `int data[10]`
↑ **Heap**	**dynamic memory** `new int[10]` `malloc(40)`
BSS and Data Segment .bss/.data	**Static/Global data** `int data[10]` (global scope)
Code .text	

lower memory
addresses
0x00FF0000

Stack Memory

A local variable is either in stack memory or CPU registers

```cpp
int x = 3; // not on stack

struct A {
    int k; // depends on where the instance of A is
};

int main() {
    int  y   = 3;        // on stack
    char z[] = "abc";    // on stack
    A    a;              // on stack (also k)
    int* ptr = new int;  // variable "ptr" is on stack
}
```

The organization of stack memory enables much higher performance. On the other hand, this memory space is limited!!
It is $\approx 8MB$ on linux by default

Stack Memory

Every object which resides in the stack is not valid outside his scope!!

```cpp
int* f() {
    int array[3] = {1, 2, 3};
    return array;
}
int* ptr = f();
cout << ptr[0]; // Illegal memory access!! ☠
```

```cpp
void g(bool x) {
    const char* str = "abc";
    if (x) {
        char xyz[] = "xyz";
        str        = xyz;
    }
    cout << str; // Illegal memory access!! ☠
}
```

new, delete

new/new[] and delete/delete[] are C++ *keywords* that
perform <u>dynamic</u> memory allocation/deallocation, and object
construction/destruction at runtime

malloc and free are C functions and they allocate and free
memory blocks (expressed in bytes)

Example:

```cpp
int* array = new int[10];  // C: (int*) malloc(10 * sizeof(int))
delete[] array;            // C: free(array)
```

`new, delete` **Advantages**

- **Language keywords**, not functions → *safer*

- **Return type**: `new` returns exact data type, while `malloc()` returns `void*`

- **Failure**: `new` throws an *exception*, while `malloc()` returns a `NULL` pointer → *it cannot be ignored*

- **Allocated bytes**: The size of the allocated memory is calculated by the compiler for `new`, while the user must take care of manually calculate the size for `malloc()`

- **Initialization**: `new` can be used to initialize an object or a set of objects

Dynamic Allocation

- **Allocate a single element**

```cpp
int* value = (int*) malloc(sizeof(int)); // C
int* value = new int;                    // C++
```

- **Allocate N elements**

```cpp
int* array = (int*) malloc(N * sizeof(int)); // C
int* array = new int[N];                      // C++
```

- **Allocate and zero-initialize N elements**

```cpp
int* array = (int*) calloc(N * sizeof(int)); // C
int* array = new int[N]();                    // C++
```

- **Allocate N structures**

```cpp
MyStruct* array = (int*) malloc(N * sizeof(MyStruct)); // C
MyStruct* array = new MyStruct[N];                      // C++
```

Dynamic Deallocation

- **Deallocate a single element**

```cpp
int* value = (int*) malloc(sizeof(int)); // C
free(value);

int* value = new int;                     // C++
delete value;
```

- **Deallocate N elements**

```cpp
int* value = (int*) malloc(N * sizeof(int)); // C
free(value);

int* value = new int[N];                     // C++
delete[] value;
```

Fundamental rules:

- Each object allocated with `new` must be deallocated with `delete`

- Each object allocated with `new[]` must be deallocated with `delete[]`

Mixing `new`, `new[]`, `malloc` with something different from their counterparts leads to *undefined behavior*

`delete` and `delete[]` applied to `NULL`/`nullptr` pointers do not produce errors (same as `free`)

Memory Leak

Memory Leak

A **memory leak** is a dynamically allocated entity in heap memory that is no longer used by the program, but still maintained overall its execution

Problems:

- Illegal memory accesses $\rightarrow$ segmentation fault
- Undefined values $\rightarrow$ segmentation fault
- Additional memory consumption

```cpp
int main() {
    int* array = new int[10];
    array       = nullptr; // memory leak!!
} // the memory can no longer be deallocated!!
```

Note: the memory leaks are especially difficult to detect in complex code and when objects are widely used

2D Memory Allocation

Easy on stack:

```cpp
int A[3][4];
```

Dynamic Memory 2D allocation/free:

```cpp
int** A = new int*[3];
for (int i = 0; i < 3; i++)
    A[i] = new int[4];

for (int i = 0; i < 3; i++)
    delete[] A[i];
delete[] A;
```

Dynamic memory 2D allocation/free C++11:

```cpp
auto A = new int[3][4];     // allocate 3 objects of size int[4]
int  n = 3;                 // dynamic value
auto B = new int[n][4];     // ok
// auto C = new int[n][n];  // compile error
delete[] A;                 // same for B, C
```

Data and BSS Segment

```
int data[]           = {1, 2}; // DATA segment memory
int big_data[1000000] = {};    // BSS segment memory
                               // (zero-initialized)

int main() {
    int A[] = {1, 2, 3}; // stack memory
}
```

Data/BSS (Block Started by Symbol) segments are larger than
stack memory (max ≈ 1GB in general) but slower

Initialization

Variable Initialization

C++03:

```cpp
int a1;          // default initialization (undefined value)
int a2(2);       // direct (or value) initialization
int a3 = 2;      // copy initialization
int a4 = 2u;     // copy initialization (implicit)
int a5 = int(2); // copy initialization
int a6 = int();  // copy initialization (zero-initialization)
int a7 = {2};    // copy list initialization
// int a8();     // a8 is a function
```

Uniform Initialization

C++11 provides the **Uniform Initialization** syntax, namely
brace-initialization or *braced-init-list*, to initialize different entities
(variables, objects, structures, etc.) in a <u>consistent</u> way:

```cpp
int b1{2};       // direct list (value) initialization
int b2{};        // direct list (value) initialization (default value)
int b3 = {};     // copy list initialization (default value)
int b4 = int{4}; // copy initialization
```

Brace Initialization

The C++11 **brace initialization** can be also used to *safely* convert arithmetic types, preventing implicit *narrowing*, i.e potential value loss. The syntax is also more concise than modern casts

```cpp
int       b4 = -1; // ok
int       b5{-1}; // ok
//unsigned b5{-1}; //   compile error

float   f1 {10e30}; // ok
//float f2{10e40}; // compile error
//------------------------------------------------------------
// FOR CONVERSION:
int       y1{x1}; // ok (only GCC, not clang)
//unsigned y2{x1}; // compile error (unsafe)

unsigned z1 = static_cast<unsigned>(-1); // ok, also z1 = -1
unsigned z2 = static_cast<unsigned>(x1); // ok, also z1 = x1
```

Structure Initialization

```cpp
struct S {
    unsigned x, y;
};
//----------------------------------------------------------
// S s0(3, 2);   // compile error
                 // The compiler searches for a constructor
S s1 = {3, 2};   // ok
S s2 = {3, -2};  // ok in C++03, but compile error in C++11

S s3 {3, 2};     // ok, also in C++11
// S s4 {3, -2}; // compile error in C++11

S f1() { // C++03
    S s5 = {3, 2};
    return s5;
}

S f2() { return {3, 2}; } // C++11
```

Stack Array Initialization

One dimension:

```cpp
int  a[3] = {1, 2, 3}; // explicit size
int  b[]  = {1, 2, 3}; // implicit size
char c[]  = "abcd";    // implicit size
int  d[3] = {1, 2};    // d[2] = 0 -> zero/default value

int  e[4] = {0}; // all values of D are initialized to 0
int  f[3] = {};  // all values of E are initialized to 0 (C++11)
int  g[3] {};    // all values of F are initialized to 0 (C++11)
```

Two dimensions:

```cpp
int a[][2] = { {1,2}, {3,4}, {5,6} }; // ok
int b[][2] = { 1, 2, 3, 4 };          // ok
// the type of "a" and "b" is an array of type int[]
// int c[][]  = ...;                  // compile error
// int d[2][] = ...;                  // compile error
```

Dynamic Initialization

C++03:

```cpp
int* a1 = new int;       // undefined
int* a2 = new int();     // zero-initialization
int* a3 = new int(4);    // allocate a single value equal to 4
int* a4 = new int[3];    // allocate 4 elements with undefined values
int* a5 = new int[4](); // allocate 4 elements zero-initialized
// int* a6 = new int[4](3); // not valid
```

C++11:

```cpp
int* b1 = new int[4]{};       // allocate 4 elements zero-initialized
int* b2 = new int[4]{1, 2}; // set first, second, zero-initialized
```

Pointers and References

Pointer

A **pointer** T* is a value referring to a location in memory

Pointer Dereferencing

Pointer **dereferencing** (*ptr) means obtaining the value stored
in at the location refereed to the pointer

Subscript Operator []

The subscript operator (ptr[]) allows accessing to the pointer
element at a given position

Deferencing:

```
int* ptr1 = new int;
*ptr1     = 4;    // deferencing (assignment)
int a     = *ptr1; // deferencing (get value)
```

Array subscript:

```
int* ptr2 = new int[10];
ptr2[2]   = 3;
int var   = ptr2[4];
```

Common error:

```
int *ptr1, ptr2; // one pointer and one integer!!
int *ptr1, *ptr2; // ok, two pointers
```

$1 + 1 \neq 2$: **Pointer Arithmetic**

Pointer syntax:

`ptr[i]` is equal to `*(ptr + i)`

Pointer arithmetic rule:

$\text{address}(\text{ptr} + i) = \text{address}(\text{ptr}) + (\text{sizeof(T)} * i)$

where `T` is the type of elements pointed by `ptr`

Example:

```
int array[4] = {1, 2, 3, 4};
cout << array[1];      // print 2
cout << *(array + 1);  // print 2
cout << array;         // print 0xFFFAFFF2
cout << array + 1;     // print 0xFFFAFFF6!!
```

`char arr[3] = "abc"`

value	address	
'a'	0x0	←arr[0]
'b'	0x1	←arr[1]
'c'	0x2	←arr[2]

`int arr[3] = {4,5,6}`

value	address	
4	0x0	←arr[0]
	0x1	
	0x2	
	0x3	
5	4	←arr[1]
	0x5	
	0x6	
	0x7	
	0x8	←arr[2]

The **address-of operator** (&) returns the address of a variable

```cpp
int  a = 3;
int* b = &a; // address-of operator,
             // 'b' is equal to the address of 'a'
a++;
cout << *b; // print 4;
```

To not confuse with **Reference syntax:** `T& var = ...`

Wild and Dangling Pointers

Wild pointer:

```cpp
int main() {
    int* ptr;   // wild pointer: Where will this pointer points?
    ...         // solution: always initialize a pointer
}
```

Dangling pointer:

```cpp
int main() {
    int* array = new int[10];
    delete[] array; // ok -> "array" now is a dangling pointer
    delete[] array; // double free or corruption!!
    // program aborted, the value of "array" is not null
}
```

Solution:

```cpp
int main() {
    int* array = new int[10];
    delete[] array; // ok -> "array" now is a dangling pointer
    array = nullptr; // no more dagling pointer
    delete[] array; // ok, no side effect
}
```

`void` **Pointer (Generic Pointer)**

Instead of declaring different types of pointer variable it is possible to declare single pointer variable which can act as any pointer types

- A `void*` can be assigned to another `void*`
- `void*` can be <u>compared</u> for equality and inequality
- A `void*` can be <u>explicitly converted</u> to another type
- Other operations would be unsafe because the compiler cannot know what kind of object is really pointed to. Consequently, other operations result in compile-time errors

```cpp
cout << (sizeof(void*) == sizeof(int*));  // print true

int array[] = { 2, 3, 4 };
void* ptr = array;
cout << *array;          // print 2
// cout << *ptr;         // compile error
cout << *((int*) ptr);   // print 2
// void* ptr2 = ptr + 2; // compile error
```

Reference

A variable **reference** `T&` is an **alias**, namely another name for an already existing variable. Both variable and variable reference can be applied to refer the value of the variable

- A pointer has its own memory address and size on the stack, reference shares the **same memory address** (with the original variable)

- References can be internally implemented as *pointers*, but the compiler treats them in a very different way

References are safer than pointers:

- References **cannot have NULL** value. You must always be able to assume that a reference is connected to a legitimate storage

- References **cannot be changed**. Once a reference is initialized to an object, it cannot be changed to refer to another object (Pointers can be pointed to another object at any time)

- References must be **initialized** when they are created (Pointers can be initialized at any time)

Reference (Examples)

Reference syntax: `T& var = ...`

```cpp
//int& a;      // compile error no initilization
//int& b = 3;  // compile error "3" is not a variable
int  c = 2;
int& d = c;    // reference. ok valid initialization
int& e = d;    // ok. the reference of a reference is a reference
d++;           // increment
e++;           // increment
cout << c;     // print 4
```

```cpp
int  a = 3;
int* b = &a;   // pointer
int* c = &a;   // pointer
b++;           // change the value of the pointer 'b'
*c++;          // change the value of 'a' (a = 4)

int& d = a;    // reference
d++;           // change the value of 'a'  (a = 5)
```

Reference vs. pointer arguments:

```
void f(int* value) {} // value may be a nullptr

void g(int& value) {} // value is never a nullptr

int a = 3;
f(&a);   // ok
f(0);    // dangerous but it works!! (but not with other numbers)
//f(a);  // compile error "a" is not a pointer

g(a);    // ok
//g(3);  // compile error "3" is not a reference of something
//g(&a); // compile error "&a" is not a reference
```

References can be use to indicate fixed size arrays:

```cpp
void f(int (&array)[3]) { // accepts only arrays of size 3
    cout << sizeof(array);
}
void g(int array[]) {
    cout << sizeof(array); // any surprise?
}

int  A[3], B[4];
int* C = A;
//-----------------------------------------------------
f(A);   // ok
// f(B); // compile error B has size 4
// f(C); // compile error C is a pointer
g(A);   // ok
g(B);   // ok
g(C);   // ok
```

Reference (Arrays) ★

```cpp
int A[4];
int (&B)[4] = A;      // ok, reference to array
int C[10][3];
int (&D)[10][3] = C;  // ok, reference to 2D array

auto c = new int[3][4];  // type is int (*)[4]
// read as "pointer to arrays of 4 int"
// int (&d)[3][4] = c;    // compile error
// int (*e)[3]    = c;    // compile error
int (*f)[4] = c;          // ok
```

```cpp
int array[4];
// &array is a pointer to an array of size 4
int size1 = (&array)[1] - array;
int size2 = *(&array + 1) - array;
cout << size1;  // print 4
cout << size2;  // print 4
```

see also www3.ntu.edu.sg/home/

Reference and `struct`

- The `dot` (`.`) operator is applied to local objects and references
- The `arrow` operator (`->`) is used with a pointer to an object

```cpp
#include <iostream>
using namespace std;
struct A {
   int x = 3;
};
int main() {
   A obj;

   A* p = &obj;    // pointer
   p->x;           // arrow syntax

   A& ref = obj;  // reference
   cout << obj.x; // dot syntax
   cout << ref.x; // dot syntax
}
```

const **and** constexpr

const Keyword

const keyword

The const keyword indicates objects never changing value after
their initialization (they must be initialized when declared)

const variables are evaluated at compile-time value if the right
expression is also evaluated at compile-time

```cpp
int size = 3;
int A[size] = {1, 2, 3}; // Technically possible (size is dynamic)
                         // But NOT approved by the C++ standard
const int SIZE = 3;
// SIZE = 4;             // compile error (SIZE is const)
int B[SIZE] = {1, 2, 3}; // ok

const int size2 = size;
int B[size2] = {1, 2, 3}; // BAD programming!! size is not const
// (some compilers allow variable size stack array -> dangerous!!)
```

Constness rules:

- `int* → const int*`

- `const int* ↛ int*`

```
int f1(const int* array) { // the values of the array cannot be
    ...                     // modified
}

int f2(int* array) {}

int*        ptr   = new int[3];
const int* c_ptr = new int[3];
f1(ptr);      // ok
f2(ptr);      // ok
f1(c_ptr);    // ok
// f2(c_ptr); // compile error
```

```
void g(const int) { // pass-by-value combined with 'const'
    ...             // note: it is not useful because the value
}                   //       is copied
```

- `int*` pointer to int
  - The value of the pointer can be modified
  - The elements refereed by the pointer can be modified

- `const int*` pointer to const int. Read as (const int)*
  - The value of the pointer can be modified
  - The elements refereed by the pointer cannot be modified

- `int *const` const pointer to int
  - The value of the pointer cannot be modified
  - The elements refereed by the pointer can be modified

- `const int *const` const pointer to const int
  - The value of the pointer cannot be modified
  - The elements refereed by the pointer cannot be modified

Note: `const int*` is equal to `int const*`

Tip: pointer types should be read from right to left

Common error: adding `const` to a pointer is not the same as adding `const` to a type alias of a pointer

```cpp
using ptr_t       = int*;
using const_ptr_t = const int*;

void f1(const int* ptr) {
// ptr[0] = 0;      // not allowed: pointer to const objects
    ptr   = nullptr; // allowed
}
void f3(const_ptr_t ptr) { // same as before
// ptr[0] = 0;      // not allowed: pointer to const objects
    ptr   = nullptr; // allowed
}
void f2(const ptr_t ptr) { // warning!!
    ptr[0] = 0;      // allowed
// ptr   = nullptr; // not allowed: const pointer to
}                    // modifiable objects
```

constexpr (function)

C++11/C++14/C++17 guarantees compile-time evaluation of an function as long as **all** its arguments are constant

- C++11: constexpr must contain exactly one `return` statement and it must not contain loops or switch
- C++14: constexpr has no restrictions

constexpr (variable)

C++11/C++14/C++17 constexpr variables are evaluated at compile-time

- const guarantees the value of a variable to be fixed overall the execution of the program
- constexpr tells the compiler that the expression results is at compile-time. constexpr value <u>implies</u> const

```
const int v1 = 3;      // compile-time evaluation
const int v2 = v1 * 2; // compile-time evaluation

int       a = 3;       // "a" is dynamic
const int v3 = a;      // run-time evaluation!!

constexpr c1 = v1;     // ok
// constexpr c2 = v3;  // compile error
```

```
constexpr int square(int value) {
    return value * value;
}

square(4); // compile-time evaluation

int a = 4; // "a" is dynamic
square(a); // run-time evaluation
```

C++17 introduces `if constexpr` feature which allows *conditionally* compiling code based on a *compile-time* value

It is an `if` statement where the branch is chosen at compile-time (similarly to the `#if` preprocessor)

```cpp
void f() {
    if constexpr (true)
        cout << "compile!";
    else
        cout << "error!"; // never compiled
}
```

`constexpr` example

```cpp
constexpr int fib(int n) {
    return (n == 0 || n == 1) ? 1 : fib(n - 1) + fib(n - 2);
}

int main() {
    if constexpr (sizeof(void*) == 8)
        return fib(5);
    else
        return fib(3);
}
```

Generated assembly code (x64 OS):

```asm
main:
    mov eax, 8
    ret
```

Explicit Type Conversion

Old style cast `(type) value`

New style cast:

- `static_cast` does compile-time (not run-time) checking of
  the types involved In many situations, this can make it the
  safest type of cast, as it provides the least room for
  accidental/unsafe conversions between various types

- `const_cast` can cast away (remove) constness or volatility

- `reinterpret_cast`
  `reinterpret_cast<T*>(v)` equal to `(T*) v`
  `reinterpret_cast<T&>(v)` equal to `*((T*) &v)`

`const_cast` and `reinterpret_cast` do not compile to any CPU
instructions

Static cast vs. old style cast:

```cpp
char a[] = {1, 2, 3, 4};
int* b   = (int*) a;              // ok
cout << b[0];                     // print 67305985 not 1!!
//int* c = static_cast<int*>(a);  // compile error unsafe conversion
```

Const cast:

```cpp
const int      a = 5;
const_cast<int>(a) = 3; // ok
```

Reinterpret cast: (bit-level conversion)

```cpp
float b = 3.0f;
// bit representation of b: 01000000010000000000000000000000
int   c = reinterpret_cast<int&>(b);
// bit representation of c: 01000000010000000000000000000000
```

Print the value of a pointer

```cpp
int* ptr = new int;
//int x1 = static_cast<size_t>(ptr);     // compile error unsafe
int  x2 = reinterpret_cast<size_t>(ptr); // ok, same size

// but
unsigned v;
//int x3 = reinterpret_cast<int>(v); // compile error
                                     // invalid conversion
int  x4 = (int) v; // ok
```

Array reshaping

```cpp
int a[3][4];
int (&b)[2][6] = reinterpret_cast<int (&)[2][6]>(a);
int (*c)[6]    = reinterpret_cast<int (*)[6]>(a);
```

Pointer Aliasing

One pointer **aliases** another when they both point to the <u>same</u> memory location

Type Punning

Type punning refers to circumvent the type system of a programming language to achieve an effect that would be difficult or impossible to achieve within the bounds of the formal language

The compiler assumes that the *strict aliasing rule is never violated*. Accessing a value using a type which is different from the original one is not allowed and it is classified as *undefined behavior*

```
// slow without optimizations. The branch breaks the pipeline
float abs(float x) {
    return (x < 0.0f) ? -x : x;
}


// optimized by hand
float abs(float x) {
    unsigned* uvalue = (unsigned*) &x;
    return (*uvalue) & 0x7FFFFFFF; // clear the last bit
}
// this is (potentially) undefined behavior!!
```

GCC warning (not clang): `-Wstrict-aliasing`

---

blog.qt.io/blog/2011/06/10/type-punning-and-strict-aliasing 45/49

`sizeof` **Operator**

sizeof operator

sizeof

The `sizeof` is a compile-time operator that determines the size, in bytes, of a variable or data type

- `sizeof` returns a value of type size_t
- `sizeof(incomplete type)` produces compile error, e.g. void
- `sizeof(bitfield member)` produces compile error
- `sizeof(anything)` never returns 0, except for array of size 0
- `sizeof(char)` always returns 1
- When applied to structures, it also takes into account padding
- When applied to a reference, the result is the size of the referenced type

```
sizeof(int);    // 4
sizeof(int*)    // 8 on a 64-bit OS
sizeof(void*)   // 8 on a 64-bit OS
sizeof(size_t)  // 8 on a 64-bit OS
```

```
int f(int[] array) {      // dangerous!!
    cout << sizeof(array);
}

int  array1[10];
int* array2 = new int[10];
cout << sizeof(array1); // print sizeof(int) * 10 = 40
cout << sizeof(array2); // print sizeof(int*) = 8  (64-bit)
f(array1);              // print 8
```

```cpp
struct B {
    int  x;
    char y;
};

struct C : B { // C extends B
    short z;
};
//--------------------------------------------------------------
sizeof(B);    // 8 : 4 + 1 (+ 3) (padding)
sizeof(C);    // 12 : sizeof(B) + 2 (+ 2) (padding)

struct A {};
sizeof(A);    // 1 : sizeof never return 0 (except for arrays)
```

```cpp
char  a;
char& b = a;
sizeof(&a);    // 8 in a 64-bit OS (pointer)
sizeof(b);     // 1, i.e. sizeof(char)
               // NOTE: a reference is not a pointer

struct A {};
A array1[10];
sizeof(array1); // 1 : array of empty structures

int array2[0];
sizeof(array2); // 0 : special case
```