

# Modern C++ Programming

## 10. TRANSLATION UNITS

---

*Federico Busato*

University of Verona, Dept. of Computer Science  
2021, v3.16



## 1 Basic Concepts

- Translation Unit
- Local and Global Scope
- Linkage

## 2 Storage Class and Duration

- Storage Duration
- Storage Class
- `static` and `extern` Keywords
- Internal/External Linkage Examples
- Linkage of `const` and `constexpr`
- Static Initialization Order Fiasco

### **3 Dealing with Multiple Translation Units**

- One Definition Rule (ODR)
- `inline` Functions/Variables
- Class in Multiple Translation Units
- Global Variable Issues

### **4 Function Template**

### **5 Class Template**

### **6 ODR Undefined Behavior and Summary**

## 7 #include Issues

- Forward Declaration
- Include Guard
- Circular Dependencies
- Common Linking Errors

## 8 Namespace

- Namespace Functions vs. Class + static Methods
- Namespace Alias
- Anonymous Namespace
- `inline` Namespace

## 9 How to Compile

- Compile Methods
- Deal with Libraries
- Build Static/Dynamic Libraries
- Find Dynamic Library Dependencies
- Analyze Object/Executable Symbols

# Basic Concepts

---

# Translation Unit

## Header File and Source File


**Header files** allow to define interfaces (.h, .hpp, .hxx), while keeping the implementation in separated **source files** (.c, .cpp, .cxx).

## Translation Unit

A **translation unit** (or *compilation unit*) is the basic unit of compilation in C++. It consists of the content of a single source file, plus the content of any header files directly or indirectly included by it

A single translation unit can be compiled into an object file, library, or executable program

# Compile Process





# Local and Global Scope

## Scope

The **scope** of a variable/function/object is the region of the code within the entity can be accessed

## Local Scope

Variables that are declared inside a function or a block are called local variables (**local scope** or **block scope**)

## Global Scope

Variables that are defined outside of all the functions and hold their value throughout the life-time of the program are global variables (**global scope** or **file scope**)

# Local and Global Scope

```
int var1;    // global scope

int f() {
    int var2; // local scope
}

struct A {
    int var3; // depends on where an instance of A is used
};

int main() {
    int var4; // local scope
}
```

# Linkage

## Linkage

**Linkage** refers to the *visibility* of symbols to the linker

## Internal Linkage

**Internal linkage** refers to symbols visible only in scope of a *single* translation unit. The same symbol name has a different memory address in distinct translation units

## External Linkage

**External linkage** refers to entities that exist *outside* a single translation unit. They are accessible and have the same *identical memory address* through the whole program, which is the combination of all translation units

## No Linkage

**No linkage** refers to symbols visible only in the local scope of declaration

# Storage Class and Duration

---

## Storage Duration

The **storage duration** (or *duration class*) determines the *duration* of a variable, namely when it is created and destroyed

Storage Duration	Allocation	Deallocation
<b>Automatic</b>	Code block start	Code end start
<b>Static</b>	Program start	Program end
<b>Dynamic</b>	Memory allocation	Memory deallocation
<b>Thread</b>	Thread start	Thread end

- **Automatic storage duration**. Scope variables (local variable). register or stack (depending on compiler, architecture, etc.).
- **Static storage duration**. The storage of an object is allocated when the program begins and deallocated when the program ends. If not explicitly initialized, it is zero-initialized
- **Dynamic storage duration**. The object is allocated and deallocated by using dynamic memory allocation functions ( `new/delete` ). If not explicitly initialized, its memory content is undefined
- **Thread storage duration** C++11. The object is allocated when the thread begins and deallocated when the thread ends. Each thread has its own instance of the object

# Storage Duration Examples

```
int          v1;      // static duration
static int   v2 = 4;  // static duration
extern int   v3;      // static duration

void f() {
    int       v4;      // automatic duration
    auto      v5 = 3;   // automatic duration
    static int v6;      // static duration
    auto array = new int[10]; // dynamic duration (allocation)
} // array, v1, v2, v3, v6 variables deallocation (from stack)
  // the memory associated with "array" is not deallocated!!

int main() {
    auto array = new int[10]; // dynamic duration (allocation)
    delete[] array;          // dynamic duration (deallocation)
}

// main end: v1, v2, v3, v6 deallocation
```

# Local static Variable

`static` *local variables* are allocated when the program begins, *initialized* when the function is called the first time, and deallocated when the program end

```
#include <iostream>
using namespace std;

int f() {
    static int val = 1;
    val++;
    return val;
}

int main() {
    cout << f(); // print 1 ("val" is initialized)
    cout << f(); // print 2
    cout << f(); // print 3
}
```



# Storage Class

## Storage Class Specifier

The **storage class** for a variable declaration is a **type specifier** that, *together with the scope*, governs its *storage duration* and *linkage*

- Only one *storage class specifier* may appear in a declaration except for `thread_local`

Storage Class	Scope	Storage Duration	Linkage
<code>auto</code> (or <code>type</code> )	Local	Automatic	No linkage
<code>static</code>	Local	static	No linkage
<code>static</code>	Global	static	Internal
<code>extern</code>	Global	static	External
<code>thread_local</code> C++11	any	Thread Local	any

Global variables have no storage class. They have static storage duration and external linkage

# Storage Class Examples

```
int          v2;      // no storage class
static      int v3 = 2; // static storage class
extern      int v4;    // external storage class
thread_local int v5;    // thread local storage class
thread_local static int v6; // thread local and static storage classes

int main() {
    int          v7;      // auto storage class
    auto         v8 = 3;   // auto storage class
    static int    v10;     // static storage class
    thread_local int v11;   // thread local and auto storage classes
    auto array = new int[10]; // auto storage class ("array" variable)
}
```

## static and extern Keywords

`static` / *anonymous namespace-included global variables or functions* are visible only within the file (*internal linkage*)

- **Non-`static`** global variables or functions with the same name in different translation units produce name collision (or name conflict)

`extern` keyword is used to declare the existence of *global variables or functions* in another translation unit (*external linkage*)

- the variable or function must be defined in a one and only one translation unit

If the same identifier within a translation unit appears with both *internal* and *external* linkage, the behavior is undefined

# Internal/External Linkage Examples

```
int      var1 = 3;  // external linkage
                // (in conflict with variables in other
                // translation units with the same name)

static int var2 = 4; // internal linkage (visible only in the
                // current translation unit)

extern int var3;     // external linkage
                // (implemented in another translation unit)

void     f1() {}    // external linkage (could conflict)

static void f2() {}  // internal linkage

namespace {          // anonymous namespace
void      f3() {}    // internal linkage
}

extern void f4();     // external linkage
                // (implemented in another translation unit)

int main() {}
```

# Linkage of const and constexpr

`const` variable at global scope implies `static`

→ *internal linkage*

`constexpr` implies `const`, which implies `static`

→ *internal linkage*

*note:* the same variable has different memory addresses on different translation units

```
const      int var1 = 3;          // internal linkage
constexpr int var2 = 2;          // internal linkage

static const      int var3 = 3; // internal linkage (redundant)
static constexpr int var4 = 2; // internal linkage (redundant)

int main() {}
```

In C++, the order in which global variables are initialized at runtime is not defined. This introduces a subtle problem called *static initialization order fiasco*.

source.cpp

```
int f() { return 3; } // run-time function

int x = f();          // run-time evaluation
```

main.cpp

```
#include <iostream>

extern int x;

int      y = x; // run-time initialized

int main() {
    std::cout << y; // print "3" or "0" depending on the linking order
}
```

source.cpp

```
constexpr int f() { return 3; } // compile-time/run-time function

constinit int x = f();           // compile-time initialized
```

main.cpp

```
#include <iostream>

constinit extern int x;          // compile-time initialized
int y = x;                       // run-time initialized

int main() {
    std::cout << y; // print "3"!!
}
```

# Linkage Summary

## No Linkage:

- Local variables, functions, classes

## Internal Linkage: (not visible to other translation units)

- **Variables:** Global `static`, `const`, `constexpr`
- **Functions:** `static`, `constexpr`
- Anonymous `namespace` content (even classes)

## External Linkage: (visible to other translation units)

- **Variables:** Global with no specifier
- **Functions:** no specifier, `extern`, `template`
- `static` class data member (that are not `inline`)



# Dealing with Multiple Translation Units


---

# One Definition Rule (ODR)

- (1) In any **(single) translation unit**, a template, type, function, or object, *cannot* have more than one definition
  - *Compiler error* otherwise
  - Any number of declarations are allowed
- (2) In the **entire program**, an object or non-inline function *cannot* have more than one definition
  - *Multiple definitions linking error* otherwise
  - Entities with *internal linkage* in different translation units are allowed, even if their names and types are the same
- (3) A template, type, or inline functions/variables, can be defined in more than one translation unit. For a given entity, each definition must be the same
  - *Undefined behavior* otherwise
  - Common case: same header included in multiple translation units


# One Definition Rule - Code Structure 1

- one header, two source files → two translation units
- *the header is included in both translation units*



## One Definition Rule - Code Structure 2

- two headers, two source files → two translation units
- one header for declarations (.hpp), and the other one for implementations (.i.hpp)
- *the header and the header implementation are included in both translation units*



\* separate header declaration and implementation is not mandatory but it could help to better organize the code

# One Definition Rule - Point (1), (2)

header.hpp:

```
void f();
```

main.cpp:

```
#include "header.hpp"
#include <iostream>
// external linkage
int      a = 1;

extern int b;
// internal linkage
static int c = 2;

int main() {
    std::cout << a; // print 1
    std::cout << b; // print 5
    std::cout << c; // print 2
    f();           // print 4
}
```

source.cpp:

```
#include "header.hpp"
#include <iostream>
// linking error, multiple definitions
// int      a = 2;

int      b = 5;    // ok
// internal linkage
static int c = 4;  // ok

void f() {         // definition
    std::cout << b; // print 5
}
```

## inline

`inline` specifier allows a function to be identically defined (not only declared) in multiple translation units

- `inline` is one of the most misunderstood features of C++
- `inline` is a hint for the linker. Without it, the linker can emit “multiple definitions” error
- `inline` variables/functions have *external linkage* (unique memory address) as standard variables/functions but cannot be exported, namely, used by other translation units
- C++17 `inline` can be also applied to variables
- It can be applied for optimization purposes only if a function has *internal linkage* (`static` or it within an `anonymous namespace`)

```
void f() {}  
inline void g() {}
```

f() :

- Cannot be defined in a header included in multiple source files
- The linker issues a “multiple definitions” error

g() :

- Can be defined in a header and included in multiple source files
- The linker removes all definitions except one
- Multiple definitions do not break the ODR rule

header.hpp:

```
inline void f() {} // the function is inline (no linking error)

template<typename T>
void g(T x) {}      // the function is a template (no linking error)

using var_t = int; // types can be defined multiple times (no linking error)
```

main.cpp:

```
#include "header.hpp"

int main() {
    f();
    g(3); // g<int> generated
}
```

source.cpp:

```
#include "header.hpp"

void h() {
    f();
    g(3); // g<int> generated
}
```



## Alternative organization:

header.hpp:

```
inline void f();    // declaration
inline int gvalue;  // declar. (C++17)

template<typename T>
void g(T x);        // declaration

using var_t = int;  // type
#include "header.i.hpp"
```

header.i.hpp:

```
void f() {}         // definition
int gvalue = 3;     // def. (C++17)

template<typename T>
void g(T x) {}      // definition
```

main.cpp:

```
#include "header.hpp"

int main() {
    f();
    g(3); // g<int> generated
}
```

source.cpp:

```
#include "header.hpp"

void h() {
    f();
    g(3); // g<int> generated
}
```

header.hpp:

```
class A {  
public:  
    void f();  
    static void g();  
private:  
    int x;  
    static int y;  
};
```

main.cpp:

```
#include "header.hpp"  
#include <iostream>  
  
int main() {  
    A a;  
    std::cout << A.x; // print 1  
    std::cout << A.y; // print 2  
}
```

source.cpp:

```
#include "header.hpp"  
  
void A::f() {}  
void A::g() {}  
  
int A::x = 1;  
int A::y = 2;
```

header.hpp:

```
struct A {  
    int x1;  
    int x2 = 3;  
    int x3 { 4 };  
  
    static int y;    // zero-init  
// static int y = 3; // compile error  
//           must be initialized out-of-class  
  
    const int z = 3; // only in C++11  
// const int z;     // compile error  
//           must be initialized  
  
    static const int w1;    // zero-init  
    static const int w2 = 4; // inline-init  
};
```

source.cpp:

```
#include "header.hpp"  
  
int      A::x1 = 1;  
int      A::y  = 2;  
const int A::w1 = 3;
```

# ODR Common Class Errors

header.hpp:

```
struct A {  
    void f() {}; // inline definition  
    void g();    // declaration  
    void h();    // declaration  
};  
  
void A::g() {} // definition!!
```

main.cpp:

```
#include "header.hpp"  
// linking error  
// multiple definitions of A::g()  
  
int main() {  
}
```

source.cpp:

```
#include "header.hpp"  
// linking error  
// multiple definitions of A::g()  
  
void A::h() {} // definition, ok
```

# Global Variable Issues

header.hpp:

```
#include <iostream>

struct A {
    A() { std::cout << "A()"; }
    ~A() { std::cout << '~A()'; }
};

// A          obj;          // linking error multiple definitions
const A       const_obj;   // "const" implies static as "constexpr"
constexpr float PHI = 3.14f;
```

source1.cpp:

```
#include "header.hpp"

void f() { std::cout << &PHI; }
// address: 0x1234ABCD

// print "A()" the first time
// print "~A()" the first time
```

source2.cpp:

```
#include "header.hpp"

void f() { std::cout << &PHI; }
// print address: 0x3820FDAC !!

// print "A()" the second time!!
// print "~A()" the second time!!
```

# Function Template

---

# Function Template - Case 1

header.hpp:

```
template<typename T>
void f(T x) {}; // declaration and definition
```

main.cpp:

```
#include "header.hpp"

int main() {
    f(3);    // call f<int>()
    f(3.3f); // call f<float>()
    f('a');  // call f<char>()
}
```

source.cpp:

```
#include "header.hpp"

void h() {
    f(3);    // call f<int>()
    f(3.3f); // call f<float>()
    f('a');  // call f<char>()
}
```

`f<int>()` , `f<float>()` , `f<char>()` are generated in both translation units

## Function Template Specialization - Case 2

header.hpp:

```
template<typename T>
void f(T x); // DECLARATION
```

main.cpp:

```
#include "header.hpp"

int main() {
    f(3);    // call f<int>()
    f(3.3f); // call f<float>()
    // f('a'); // linking error
} // specialization does not
   // exist
```

source.cpp:

```
#include "header.hpp"

template<typename T>
void f(T x) {} // DEFINITION

// template SPECIALIZATION
template void f<int>(int);
template void f<float>(float);
// any explicit instance is also
// fine, e.g. f<int>(3)
```



## Function Template Specialization - Case 3a

header.hpp:

```
template<typename T>
void f(T x) {} // DECLARATION and DEFINITION

template<>
void f<int>(); // SPECIALIZATION declaration
               // inform that the specialization exists in
               // another translation unit
```

main.cpp:

```
#include "header.hpp"

int main() {
    f<int>(); // ok
    // f<char>(); // linking error
}
```

source.cpp:

```
#include "header.hpp"

// SPECIALIZATION
template<>
void f<int>(int x) {}
```

# Function Template Specialization - Case 3b

C++11

header.hpp:

```
template<typename T>
void f(T x) {} // DECLARATION and DEFINITION
```

main.cpp:

```
#include "header.hpp"

extern template void f<int>();
// f<int>() is not generated
// by the compiler in this
// translation unit

int main() {
}
```

source.cpp:

```
#include "header.hpp"

// SPECIALIZATION
template<>
void f<int>(int x) {}

// or any instantiation of
// f<int>()
// or
// template void f<int>(int);
```

header.hpp:

```
template<typename T>
void f();

// template<>           // linking error
// void f<int>() {}      // multiple definitions -> included twice
                        // full specializations are standard functions
                        // it can be solved by adding "inline"
```

main.cpp:

```
#include "header.hpp"

int main() {}
```

source.cpp:

```
#include "header.hpp"

// some code
```

header.hpp:

```
template<typename T>
void f();
```

main.cpp:

```
#include "header.hpp"

int main() {
// f<int>(); // linking error
//           // f<int> undefined
    f<char>(); // ok
}
```

source.cpp:

```
#include "header.hpp"

template<typename T>
void f() {} // valid in this
//           // translation unit

void g() {
    f<char>(); // generates f<char>
}
```

# Class Template

---

header.hpp:

```
template<typename T>
struct A {
    T    x = 3;
    void f() {}; // "inline" definition
};
```

main.cpp:

```
#include "header.hpp"

int main() {
    A<int>    a1; // ok
    A<float>  a2; // ok
    A<char>   a3; // ok
}
```

source.cpp:

```
#include "header.hpp"

int g() {
    A<int>    a1; // ok
    A<float>  a2; // ok
    A<char>   a3; // ok
}
```

header.hpp:

```
template<typename T>
struct A {
    T    x;
    void f(); // declaration
};
#include "header.i.hpp"
```

header.i.hpp:

```
template<typename T>
T A<T>::x = 3; // initialization

template<typename T>
void A<T>::f() {} // definition
```

main.cpp:

```
#include "header.hpp"

int main() {
    A<int>    a1; // ok
    A<float>  a2; // ok
    A<char>   a3; // ok
}
```

source.cpp:

```
#include "header.hpp"

int g() {
    A<int>    a1; // ok
    A<float>  a2; // ok
    A<char>   a3; // ok
}
```

# Class Template Specialization - Case 1

header.hpp:

```
template<typename T>
struct A {
    T    x;
    void f(); // declaration
};
```

main.cpp:

```
#include "header.hpp"

int main() {
    A<int>  a1; // ok
    // A<char> a2; // linking error
}
```

source.cpp:

```
#include "header.hpp"

template<typename T>
int A<T>::x = 3; // initialization

template<typename T>
void A<T>::f() {} // definition

// template specialization
template class A<int>;
```



# Class Template Specialization - Case 2

C++11

header.hpp:

```
template<typename T>
struct A {
    T    x;
    void f();
};
```

source.cpp:

```
#include "header.hpp"

extern template class A<int>;
// A<int> is not generated by the
// compiler in this translation unit
int main() {
    A<int> a;
}
```

source.cpp:

```
#include "header.hpp"

// template specialization
template class A<int>;

// or any instantiation of
// A<int>, e.g.  A<int> a;
```

# ODR Undefined Behavior and Summary

---

# Undefined Behavior - inline Function

main.cpp:

```
#include <iostream>
inline int f() { return 3; }

void g();

int main() {
    std::cout << f(); // print 3
    std::cout << g(); // print 3!!
}                      // not 5
```

source.cpp:

```
// same signature and inline
inline int f() { return 5; }

int g() { return f(); }
```

The linker can *arbitrary* choose one of the two definitions of `f()` . With `-O3` , the compiler could *inline* `f()` in `g()` , so now `g()` return `5`

This issue is easy to detect in trivial examples but hard to find in large codebase

*Solution:* `static` or `anonymous namespace`

# Undefined Behavior - Member Function

main.cpp:

```
#include <iostream>

struct A {
    int f() { return 3; }
};

int g();
```

main.cpp:

```
#include "header.hpp"
using namespace std;

int main() {
    A a;
    cout << a.f(); // print 3
    cout << g();  // print 3!!
}
```

source.cpp:

```
struct A {
    int f() { return 5; }
};

int g() {
    A<int> a;
    return a.f();
}
```

# Undefined Behavior - Function Template

main.cpp:

```
template<typename T>
int f() {
    return 3;
}

int g();
```

source1.cpp:

```
#include "header.hpp"
using namespace std;

int main() {
    cout << f<int>(); // print 3
    cout << g();      // print 3!!
}
```

source2.cpp:

```
template<typename T>
int f() {
    return 5;
}

int g() {
    return f<int>();
}
```

# Undefined Behavior

Other ODR violations are even harder (if not impossible) to find, e.g. Diagnosing Hidden ODR Violations in Visual C++

Some tools for partially detecting ODR violations:

- `-detect-odr-violations` flag for gold/llvm linker
- `-Wodr -flto` flag for GCC
- Clang address sanitizer + `ASAN_OPTIONS=detect_odr_violation=2`  
(link)

Another solution could be include all files in a single translation unit

- **header:** declaration of
  - functions, structures, classes, types, alias
  - `template` function, structs, classes
  - `extern` variables, functions
  - global `const/constexpr` variables
- **header implementation:** definition of
  - `inline` functions, variables
  - `template` functions, classes
- **source file:** definition of
  - functions
  - `template` full specialization functions, classes
  - `static` global variables (+ declaration)
  - `extern` variables, functions

# #include Issues

---



# Forward Declaration

**Forward declaration** is a declaration of an identifier for which a complete definition has not yet given

“*forward*” means that an entity is declared before it is used

main.cpp:

```
void f(); // function forward declaration

class A; // class forward declaration

int main() {
    f(); // ok, f() is a function and not a variable
    // A a; // compiler error no definition (incomplete type)
           // e.g. the compiler is not able to deduce the size of A
    A* a; // ok
}
```

source.cpp:

```
void f() {} // definition of f()
class A {}; // definition of A()
```

# Forward Declaration vs. `#include`

## Advantages:

- Forward declarations can save compile time as `#include` forces the compiler to open more files and process more input
- Forward declarations can save on unnecessary recompilation. `#include` can force your code to be recompiled more often, due to unrelated changes in the header

## Disadvantages:

- Forward declarations can hide a dependency, allowing user code to skip necessary recompilation when headers change
- A forward declaration may be broken by subsequent changes to the library
- Forward declaring multiple symbols from a header can be more verbose than simply `#including` the header

The `include guard` avoids the problem of multiple inclusions of a header file in a translation unit

`header.hpp`:

```
#ifndef HEADER_HPP // include guard
#define HEADER_HPP

... many lines of code ...


#endif // HEADER_HPP
```

`#pragma once` preprocessor directive is an alternative to the `include guard` to force current file to be included only once in a translation unit

- `#pragma once` is less portable but less verbose and compile faster than the `include guard`

The `include guard`/`#pragma once` should be used in every header file

Common case:



header\_A.hpp:

```
#pragma once // prevents "multiple definitions" linking error
```

```
struct A {  
};
```

header\_B.hpp:

```
#include "header_A.hpp" // included here
```

```
struct B {  
    A a;  
};
```


main.cpp:

```
#include "header_A.hpp" // .. and included here
```

```
#include "header_B.hpp"
```

```
int main() {  
    A a; // ok, here we need "header_A.hpp"  
    B b; // ok, here we need "header_B.hpp"  
}
```

A **circular dependency** is a relation between two or more modules which either directly or indirectly depend on each other to function properly



Circular dependencies can be solved by using forward declaration, or better, by rethinking the project organization

header\_A.hpp:

```
#pragma once
#include "header_B.hpp"
class A {
    B* b;
};
```

header\_B.hpp:

```
#pragma once
#include "header_C.hpp"
class B {
    C* c;
};
```

header\_C.hpp:

```
#pragma once
#include "header_A.hpp"
class C { // compile error "header_A" already included by "main.cpp"
    A* a; // the compiler cannot view the "class C"
};
```

header\_A.hpp:

```
#pragma once
class B;    // forward declaration
           // note: does not include "header_B.hpp"

class A {
    B* b;
};
```

header\_B.hpp:

```
#pragma once
class C;    // forward declaration
class B {
    C* c;
};
```

header\_C.hpp:

```
#pragma once
class A;    // forward declaration
class C {
    A* a;
};
```



# Common Linking Errors

Very common *linking* errors:

- **undefined reference**

*Solutions:*

- Check if the right headers and sources are included
- Break circular dependencies

- **multiple definitions**

*Solutions:*

- **inline** function, variable definition or **extern** declaration
- Add include guard/ **#pragma once** to header files
- Place template definition in header file and full specialization in source files

# Namespace

---

The problem: Named entities, such as variables, functions, and compound types declared outside any block has *global scope*, meaning that its name is valid anywhere in the code

***Namespaces*** allow to group named entities that otherwise would have global scope into narrower scopes, giving them ***namespace scope*** (where *std* stands for “standard”)

Namespaces provide a method for preventing name conflicts in large projects. Symbols declared inside a namespace block are placed in a named scope that prevents them from being mistaken for identically-named symbols in other scopes

# Namespace Functions vs. Class + static Methods

## Namespace functions:

- Namespace can be extended anywhere (without control)
- Namespace specifier can be avoided with the keyword `using`

## Class + `static` methods:

- Can interact only with static data members
- `struct/class` cannot be extended outside their declarations

→ `static` methods should define operations strictly related to an object state (*statefull*)

→ otherwise `namespace` should be preferred (*stateless*)

# Defining a Namespace

```
#include <iostream>

namespace ns1 {
void f() {
    std::cout << "ns1" << std::endl;
}
} // namespace ns1

namespace ns2 {
void f() {
    std::cout << "ns2" << std::endl;
}
} // namespace ns2

int main () {
    ns1::f(); // print "ns1"
    ns2::f(); // print "ns2"
    // f();    // compile error f() is not visible
}
```

# Namespace Conflicts

```
#include <iostream>

void f() { std::cout << "global" << endl; }

namespace ns1 {
    void f() { std::cout << "ns1::f()" << endl; }
    void g() { std::cout << "ns1::g()" << endl; }
}

int main () {
    f();          // ok, print "global"
    // g();       // compile error g() is not visible

    using namespace ns1; // expand "ns1" in this scope (from this line)
    // f();        // compile error ambiguous function name
    ::f();         // ok, print "global"
    ns1::f();      // ok, print "ns1::f()"
    g();           // ok, print "ns1::g()", only one choice
}
```

# Nested Namespaces

```
#include <iostream>
using namespace std;
namespace ns1 {
    void f() { cout << "ns1::f()" << endl; }
namespace ns2 {
    void f() { cout << "ns1::ns2::f()" << endl; }
} // namespace ns1
} // namespace ns2

namespace ns1 { // the same namespace can be declared multiple times,
namespace ns2 { // and extended in multiple files
    void g() {}
}
}
```

C++17 allows nested namespace definitions:

```
namespace ns1::ns2 {
    void h()
}
```

# Namespace Precedence

```
#include <iostream>

void f() { std::cout << "global"; }

namespace ns1 {
    void f() { std::cout << "ns1"; }

    void g() { f(); }
}

void g() { f(); }

int main() {
    g();          // "global"
    ns1::g();     // "ns1"
}
```



# Namespace Alias

**Namespace alias** allows declaring an alternate name for an existing namespace

```
namespace very_very_long_namespace {  
    void g() {}  
}  
  
int main() {  
    namespace ns = very_very_long_namespace; // namespace alias  
    ns::g();  
}
```

# Anonymous Namespace

A namespace with no identifier before an opening brace produces an **unnamed/anonymous namespace**

Entities inside an anonymous namespace are used for declaring unique identifiers, visible in the same source file

## Anonymous namespaces vs. static global entities

- Anonymous namespaces allow *type declarations*, and they are *less verbose*

main.cpp

```
#include <iostream>
namespace { // anonymous
    void f() { std::cout << "main"; }
}           // internal linkage

int main() {
    f();    // ok, print "main"
}
```

source.cpp

```
#include <iostream>
namespace { // anonymous
    void f() { std::cout << "source"; }
}

int g() {
    f();    // ok, print "source"
}
```

## inline Namespace

**inline namespace** is a concept similar to library versioning. It is a mechanism that makes a nested namespace look and act as if all its declarations were in the surrounding namespace

```
namespace ns1 {  
    inline namespace V99 {  
        void f(int) {}    // most recent version  
    }  
    namespace V98 {  
        void f(int) {}  
    }  
}  
  
using namespace ns1;  
  
int main() {  
    V98::f(1);    // call V98  
    V99::f(1);    // call V99  
    f(1);         // call default version (V99)  
}
```

# How to Compile

---

# Compile Methods

## Method 1

Compile all files together (naive):

```
g++ -I include/ main.cpp source.cpp -o main.x
```

**-I**: Specify the **include path** to the compiler. It can be used multiple times

## Method 2

Compile each *translation unit* in a file object:

```
g++ -c -I include/ source.cpp -o source.o
```

```
g++ -c -I include/ main.cpp -o main.o
```

Link all file objects:

```
g++ main.o source.o -o main.x
```

A **library** is a package of code that is meant to be reused by many programs

A **static library** is a set of object files that are directly linked into the final executable. If a program is compiled with a static library, all the functionality of the static library becomes part of final executable

- A static library cannot be modified without re-link the final executable
- Increase the size of the final executable
- The linker can optimize the final executable (*link time optimization*)

Given the static library `my_lib`, the corresponding file is:

- Linux: `libmy_lib.a`
- Windows: `my_lib.lib`

A **dynamic library**, also called a **shared library**, consists of routines that are loaded into the application at run-time. If a program is compiled with a dynamic library, the library does not become part of final executable. It remains as a separate unit

- A dynamic library can be modified without re-link
- Dynamic library functions are called outside the executable
- Neither the linker, nor the compiler can optimize the code between shared libraries and the final executable
- The environment variables must be set to the right shared library path, otherwise the application crashes at the beginning

Given the shared library `my_lib`, the corresponding file is:

- Linux: `libmy_lib.so`
- Windows: `my_lib.dll` + `my_lib.lib`

## Deal with Libraries

Specify the **library path** (path where search for static/dynamic libraries) to the compiler: `g++ -L<library_path> main.cpp -o main`

`-L` can be used multiple times ( `/LIBPATH` on Windows)

Specify the **library name** (e.g. `liblibrary.a`) to the compiler:

```
g++ -llibrary main.cpp -o main
```

The full path on Windows instead



# Deal with Libraries

## Linux/Unix environmental variables:

- `LIBRARY_PATH` Specify the directories where search for *static* libraries `.a` at *compile-time*. Used by the compiler
- `LD_LIBRARY_PATH` Specify the directories where search for *dynamic/shared* libraries `.so` at *run-time*. Used by the program

## Windows environmental variables:

- `LIBPATH` Specify the directories where search for *static* libraries `.lib` at *compile-time*. Used by the compiler
- `PATH` Specify the directories where search for *dynamic/shared* libraries `.dll` at *run-time*. Used by the program

# Build Static/Dynamic Libraries

## Static Library Creation

- Create object files for each translation unit (.cpp)
- Create the static library by using the **archiver** (**ar**) linux utility

```
g++ source1.c -c source1.o
g++ source2.c -c source2.o
ar rvs libmystaticlib.a source1.o source2.o
```

## Dynamic Library Creation

- Create object files for each translation unit (.cpp). Since library cannot store code at fixed addresses the compile must generate *position independent code*
- Create the dynamic library

```
g++ source1.c -c source1.o -fPIC
g++ source2.c -c source2.o -fPIC
g++ source1.o source2.o -shared -o libmydynamiclib.so
```

## Find Dynamic Library Dependencies

The `ldd` utility shows the shared objects (shared libraries) required by a program or other shared objects

```
$ ldd /bin/ls
linux-vdso.so.1 (0x00007ffcc3563000)
libselinux.so.1 => /lib64/libselinux.so.1 (0x00007f87e5459000)
libcap.so.2 => /lib64/libcap.so.2 (0x00007f87e5254000)
libc.so.6 => /lib64/libc.so.6 (0x00007f87e4e92000)
libpcre.so.1 => /lib64/libpcre.so.1 (0x00007f87e4c22000)
libdl.so.2 => /lib64/libdl.so.2 (0x00007f87e4a1e000)
/lib64/ld-linux-x86-64.so.2 (0x00005574bf12e000)
libattr.so.1 => /lib64/libattr.so.1 (0x00007f87e4817000)
libpthread.so.0 => /lib64/libpthread.so.0 (0x00007f87e45fa000)
```

The `nm` utility provides information on the symbols being used in an object file or executable file

```
$ nm -D -C something.so
```

```
  w __gmon_start__
```

```
  D __libc_start_main
```

```
  D free
```

```
  D malloc
```

```
  D printf
```

```
# -C: Decode low-level symbol names
```

```
# -D: accepts a dynamic library
```

`readelf` displays information about ELF format object files

```
$ readelf --symbols something.so | c++filt
... OBJECT LOCAL DEFAULT 17 __frame_dummy_init_array_
... FILE LOCAL DEFAULT ABS prog.cpp
... OBJECT LOCAL DEFAULT 14 CC1
... OBJECT LOCAL DEFAULT 14 CC2
... FUNC LOCAL DEFAULT 12 g()
```

*# --symbols: display symbol table*

`objdump` displays information about object files

```
$ objdump -t -C something.so | c++filt
... df *ABS*      ...  prog.cpp
...  0 .rodata    ...  CC1
...  0 .rodata    ...  CC2
...  F .text      ...  g()
...  0 .rodata    ...  (anonymous namespace)::CC3
...  0 .rodata    ...  (anonymous namespace)::CC4
...  F .text      ...  (anonymous namespace)::h()
...  F .text      ...  (anonymous namespace)::B::j1()
...  F .text      ...  (anonymous namespace)::B::j2()
```

*# --t: display symbols*

*# -C: Decode low-level symbol names*

## References and Additional Material

- 20 ABI (Application Binary Interface) breaking changes every C++ developer should know
- Policies/Binary Compatibility Issues With C++
- 10 differences between static and dynamic libraries every C++ developer should know