

Modern C++ Programming

10. TRANSLATION UNITS

Federico Busato

University of Verona, Dept. of Computer Science
2021, v3.11



1 Basic Concepts

- Translation Unit
- Local and Global Scopes

2 Linkage

- `static` and `extern` Keywords
- Internal/External Linkage Example
- Linkage of `const` and `constexpr`
- Linkage of `inline` Functions/Variables

3 Storage Class and Duration

- Storage Class
- Storage Duration

4 Dealing with Multiple Translation Units

- One Definition Rule (ODR)
- Class in Multiple Translation Units
- Global Constants

5 Function Template

6 Class Template

7 Undefined Behavior and Summary

8 #include Issues

- Forward Declaration
- Include Guard
- Circular Dependencies
- Common Linking Errors

9 Namespace

- Namespace Functions vs. Class + static Methods
- Namespace Alias
- Anonymous Namespace
- `inline` Namespace

10 How to Compile

- Compile Methods
- Compile with Libraries
- Build Static/Dynamic Libraries
- Find Dynamic Library Dependencies
- Analyze Object/Executable Symbols

Basic Concepts

Translation Unit

Header File and Source File

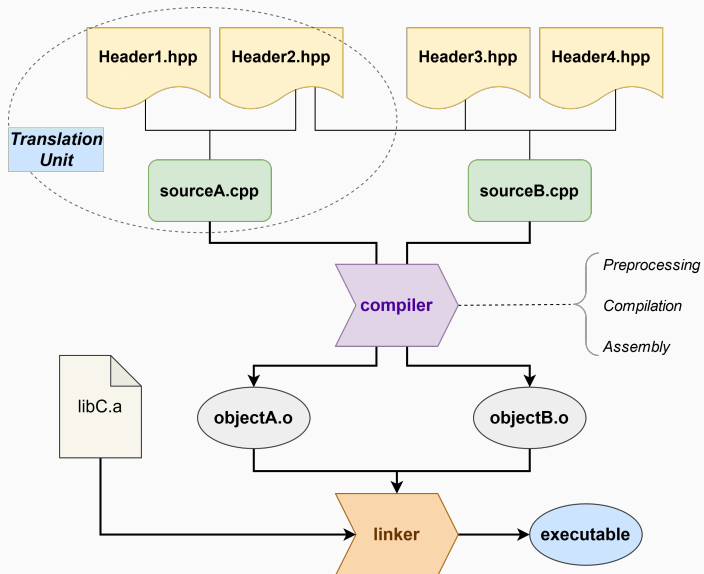
Header files allow to define interfaces (.h, .hpp, .hxx), while keeping the implementation in separated **source files** (.c, .cpp, .cxx).

Translation Unit

A **translation unit** (or *compilation unit*) is the basic unit of compilation in C++. It consists of the contents of a single source file, plus the contents of any header files directly or indirectly included by it

A single translation unit can be compiled into an object file, library, or executable program

Compile Process



Local and Global Scopes

Scope

The **scope** of a variable/function/object is the region of the code within the entity can be accessed

Local Scope

Variables that are declared inside a function or a block are called local variables (**local scope** or **block scope**)

Global Scope

Variables that are defined outside of all the functions and hold their value throughout the life-time of the program are global variables (**global scope** or **file scope**)

Local and Global Scopes

```
int var1;    // global scope

int f() {
    int var2; // local scope
}

struct A {
    int var3; // depends on where an instance of A is used
};

int main() {
    int var4; // local scope
}
```

Linkage

Linkage

Linkage

Linkage refers to the *visibility* of symbols to the linker

Internal Linkage

Internal linkage refers to symbols visible only in scope of a *single* translation unit

External Linkage

External linkage refers to entities that exist *outside* a single translation unit. They are accessible through the whole program, which is the combination of all translation units

No Linkage

No linkage refers to symbols visible only in the local scope of declaration

static and extern Keywords

`static` / *anonymous namespace-included global variables or functions* are visible only within the file (*internal linkage*)

- **Non-`static`** global variables or functions with the same name in different translation units produce name collision (or name conflict)

`extern` keyword is used to declare the existence of *global variables or functions* in another translation unit (*external linkage*)

- the variable or function must be defined in a one and only one translation unit

If, within a translation unit, the same identifier appears with both *internal* and *external* linkage, the behavior is undefined

static Variable Example

To not confuse with symbol visibility

```
#include <iostream>
using namespace std;

int f() {
    static int val = 1;  // static
    val++;
    return val;
}

int main() {
    cout << f(); // print 1
    cout << f(); // print 2
    cout << f(); // print 3
}
```

Internal/External Linkage Example

```
int      var1 = 3;  // external linkage
                // (in conflict with variables in other
                // translation units with the same name)

static int var2 = 4; // internal linkage (visible only in the
                // current translation unit)

extern int var3;     // external linkage
                // (implemented in another translation unit)

void     f1() {}    // external linkage (could conflict)

static void f2() {}  // internal linkage

namespace {          // anonymous namespace
void      f3() {}    // internal linkage
}

extern void f4();     // external linkage
                // (implemented in another translation unit)

int main() {}
```

Linkage of const and constexpr

`const` variable at global scope implies `static`

→ *internal linkage*

`constexpr` implies `const`, which implies `static`

→ *internal linkage*

note: the same variable has different memory addresses on different translation units

```
const      int var1 = 3;      // internal linkage
constexpr int var2 = 2;      // internal linkage

static const      int var3 = 3; // internal linkage (redundant)
static constexpr int var4 = 2; // internal linkage (redundant)

int main() {}
```

`inline`

`inline` specifier allows a function to be defined identically (not only declared) in multiple translation units (source file + headers)

- `inline` is one of the most misunderstood features of C++
- `inline` is a hint for the linker. Without it, the linker can emit “multiple definitions” error
- It can be applied for optimization purposes only if the function has *internal* linkage (`static` or inside an `anonymous namespace`)
- C++17 `inline` can be also applied to variables

Multiple definitions of the same `inline`-declared function with *external linkage* do not break the ODR rule (allowed behavior)

note: the same `inline` variable with *external linkage* has the same memory address on different translation units

```
inline      int var1 = 3;  // external linkage (C++17)
static inline int var2 = 3; // internal linkage (C++17)

inline      void f() {}   // external linkage
static      void g1() {}  // internal linkage
static inline void g2() {} // internal linkage
namespace {
    inline void g3() {}    // internal linkage
} // anonymous namespace -> same as static
```

`f()` :

- Can be defined in a header and included in multiple source files
- The linker removes all definitions except one
- Declaring `void f();` in a file that does not include the header is still valid because the function has *external* linkage

`g1(), g2(), g3()` :

- Can be defined in a header included in multiple source files
- The compiler replicates the code in each translation unit (the linker does not see these functions)
- Declaring `void g1();` in a file that does not include the header is no more valid because the function has *internal* linkage

No Linkage:

- Local **Variables, Functions, Classes**

Internal Linkage:

- **Variables:** Global `static`, `const`, `constexpr`
- **Functions*:** `static`, `constexpr`
- Anonymous `namespace` content

External Linkage:

- **Variables:** Global
- **Functions*:** `extern`, `template`
- `static` class data member (that are not `inline`)

* Windows (MSVC) treats function visibility in a different way gcc.gnu.org/wiki/Visibility

Storage Class and Duration

Storage Class

Storage Class Specifier

The **storage class** for a variable declaration is a type **specifier** that, *together with the scope*, governs its *storage duration* and its *linkage*

- Only one storage class specifier may appear in a declaration except that `thread_local`

Storage Class	Scope	Storage Duration	Linkage
<code>register</code> [†]	Local	Automatic	No linkage
<code>auto</code> *	Local	Automatic	No linkage
<code>auto</code> *	Global	Automatic	External
<code>static</code>	Local	static	No linkage
<code>static</code>	Global	static	Internal
<code>extern</code>	Global	static	External
<code>thread_local</code> *	any	Thread Local	any

Storage Class Examples

```
int          v2;      // automatic storage
static      int v3 = 2; // static storage (global)
extern      int v4;    // external storage
thread_local int v5;    // thread local storage
thread_local static int v6; // thread local storage

int main() {
    int          v7;      // automatic storage
    auto         v8 = 3;   // automatic storage
    register int  v9;      // register storage (deprecated!)
    static int    v10;     // static storage (local)
    thread_local int v11;   // thread local storage
    auto array = new int[10]; // automatic storage
                                // (array variable)
}
```

Storage Duration

The **storage duration** (or *duration class*) determines the *duration* of a variable, namely when it is created and destroyed

Storage Duration	Allocation	Deallocation
Automatic	Code block start	Code end start
Static	Program start	Program end
Dynamic	Memory allocation	Memory deallocation
Thread	Thread start	Thread end

- **Automatic storage duration**. Scope variables (local variable). register or stack (depending on compiler, architecture, etc.). `register` hints to the compiler to place the object in the processor registers (deprecated in C++11)
- **Static storage duration**. The storage for the object is allocated when the program begins and deallocated when the program ends (`static` keyword at local or global scope)
- **Thread storage duration** C++11. The object is allocated when the thread begins and deallocated when the thread ends. Each thread has its own instance of the object. (`thread_local` can appear together with `static` or `extern`)
- **Dynamic storage duration**. The object is allocated and deallocated per request by using dynamic memory allocation functions (`new/delete`)

Storage Duration Examples

```
int          v1;      // static duration
static int   v2 = 4;  // static duration
extern int   v3;      // static duration

void f() {
    int       v4;      // automatic duration
    auto      v5 = 3;   // automatic duration
    static int v6;      // static duration
    auto array = new int[10]; // dynamic duration (allocation)
} // array, v1, v2, v3, v6 variables deallocation (from stack)
   // the memory associated with "array" is not deallocated!!

int main() {
    auto array = new int[10]; // dynamic duration (allocation)
    delete[] array;           // dynamic duration (deallocation)
}

// main end: v1, v2, v3, v6 deallocation
```


Dealing with Multiple Translation Units

One Definition Rule (ODR)

- (1) In any **(single) translation unit**, a template, type, function, or object, *cannot* have more than one definition
 - *Compiler error* otherwise
 - Any number of declarations are allowed
- (2) In the **entire program**, an object or non-inline function *cannot* have more than one definition
 - *Multiple definitions linking error* otherwise
 - Entities with *internal linkage* in different translation units are allowed, even if their names and types are the same
- (3) A template, type, or inline functions/variables, can be defined in more than one translation unit. For a given entity, each definition must be the same
 - *Undefined behavior* otherwise
 - Common case: same header included in multiple translation units


One Definition Rule - Code Structure 1

- one header, two source files → two translation units
- *the header is included in both translation units*



One Definition Rule - Code Structure 2

- two headers, two source files → two translation units
- one header for declarations (.hpp), and the other one for implementations (.i.hpp)
- *the header and the header implementation are included in both translation units*



* separate header declaration and implementation is not mandatory but it could help to better organize the code

One Definition Rule (Example, points (1), (2))

header.hpp:

```
void f();
```

main.cpp:

```
#include "header.hpp"
#include <iostream>
// external linkage
int      a = 1;

extern int b;
// internal linkage
static int c = 2;

int main() {
    std::cout << a; // print 1
    std::cout << b; // print 5
    std::cout << c; // print 2
    f();           // print 4
}
```

source.cpp:

```
#include "header.hpp"
#include <iostream>
// linking error, multiple definitions
// int      a = 2;

int      b = 5;    // ok
// internal linkage
static int c = 4;  // ok

void f() {         // definition
    std::cout << b; // print 5
}
```

header.hpp:

```
inline void f() {} // the function is inline (no linking error)

template<typename T>
void g(T x) {}      // the function is a template (no linking error)

using var_t = int; // types can be defined multiple times (no linking error)
```

main.cpp:

```
#include "header.hpp"

int main() {
    f();
    g(3); // g<int> generated
}
```

source.cpp:

```
#include "header.hpp"

void h() {
    f();
    g(3); // g<int> generated
}
```

Alternative organization:

header.hpp:

```
inline void f();    // declaration
inline int gvalue;  // declar. (C++17)

template<typename T>
void g(T x);        // declaration

using var_t = int;  // type
#include "header.i.hpp"
```

header.i.hpp:

```
void f() {}         // definition
int gvalue = 3;     // def. (C++17)

template<typename T>
void g(T x) {}      // definition
```

main.cpp:

```
#include "header.hpp"

int main() {
    f();
    g(3); // g<int> generated
}
```

source.cpp:

```
#include "header.hpp"

void h() {
    f();
    g(3); // g<int> generated
}
```


header.hpp:

```
class A {  
public:  
    void f();  
    static void g();  
private:  
    int x;  
    static int y;  
};
```

main.cpp:

```
#include "header.hpp"  
#include <iostream>  
  
int main() {  
    A a;  
    std::cout << A.x; // print 1  
    std::cout << A.y; // print 2  
}
```

source.cpp:

```
#include "header.hpp"  
  
void A::f() {}  
void A::g() {}  
  
int A::x = 1;  
int A::y = 2;
```

header.hpp:

```
struct A {  
    int x1;  
    int x2 = 3;  
    int x3 { 4 };  
  
    static int y;    // zero-init  
// static int y = 3; // compile error  
//           must be initialized out-of-class  
  
    const int z = 3; // only in C++11  
// const int z;     // compile error  
//           must be initialized  
  
    static const int w1;    // zero-init  
    static const int w2 = 4; // inline-init  
};
```

source.cpp:

```
#include "header.hpp"  
  
int      A::x1 = 1;  
int      A::y  = 2;  
const int A::w1 = 3;
```

ODR Common Class Errors

header.hpp:

```
struct A {  
    void f() {}; // inline definition  
    void g();    // declaration  
    void h();    // declaration  
};  
  
void A::g() {} // definition!!
```

main.cpp:

```
#include "header.hpp"  
// linking error  
// multiple definitions of A::g()  
  
int main() {  
}
```

source.cpp:

```
#include "header.hpp"  
// linking error  
// multiple definitions of A::g()  
  
void A::h() {} // definition, ok
```

Global Constants

header.hpp:

```
#include <iostream>

struct A {
    A() { std::cout << "A()"; }
    ~A() { std::cout << '~A()'; }
};

// A          obj;          // linking error multiple definitions
const A       const_obj;   // "const" implies static as "constexpr"
constexpr float PHI = 3.14f;
```

source1.cpp:

```
#include "header.hpp"

void f() { std::cout << &PHI; }
// address: 0x1234ABCD

// print "A()" the first time
// print "~A()" the first time
```

source2.cpp:

```
#include "header.hpp"

void f() { std::cout << &PHI; }
// print address: 0x3820FDAC !!

// print "A()" the second time!!
// print "~A()" the second time!!
```

Function Template

Function Template - Case 1

header.hpp:

```
template<typename T>
void f(T x) {}; // declaration and definition
```

main.cpp:

```
#include "header.hpp"

int main() {
    f(3);    // call f<int>()
    f(3.3f); // call f<float>()
    f('a');  // call f<char>()
}
```

source.cpp:

```
#include "header.hpp"

void h() {
    f(3);    // call f<int>()
    f(3.3f); // call f<float>()
    f('a');  // call f<char>()
}
```

`f<int>()` , `f<float>()` , `f<char>()` are generated in both translation units

Function Template Specialization - Case 2

header.hpp:

```
template<typename T>
void f(T x); // DECLARATION
```

main.cpp:

```
#include "header.hpp"

int main() {
    f(3);    // call f<int>()
    f(3.3f); // call f<float>()
    // f('a'); // linking error
} // specialization does not
   // exist
```

source.cpp:

```
#include "header.hpp"

template<typename T>
void f(T x) {} // DEFINITION

// template SPECIALIZATION
template void f<int>(int);
template void f<float>(float);
// any explicit instance is also
// fine, e.g. f<int>(3)
```

Function Template Specialization - Case 3a

header.hpp:

```
template<typename T>
void f(T x) {} // DECLARATION and DEFINITION

template<>
void f<int>(); // SPECIALIZATION declaration
               // inform that the specialization exists in
               // another translation unit
```

main.cpp:

```
#include "header.hpp"

int main() {
    f<int>(); // ok
    // f<char>(); // linking error
}
```

source.cpp:

```
#include "header.hpp"

// SPECIALIZATION
template<>
void f<int>(int x) {}
```


Function Template Specialization - Case 3b

C++11

header.hpp:

```
template<typename T>
void f(T x) {} // DECLARATION and DEFINITION
```

main.cpp:

```
#include "header.hpp"

extern template void f<int>();
// f<int>() is not generated
// by the compiler in this
// translation unit

int main() {
}
```

source.cpp:

```
#include "header.hpp"

// SPECIALIZATION
template<>
void f<int>(int x) {}

// or any instantiation of
// f<int>()
// or
// template void f<int>(int);
```

header.hpp:

```
template<typename T>
void f();

// template<>           // linking error
// void f<int>() {}      // multiple definitions -> included twice
                        // full specializations are standard functions
                        // it can be solved by adding "inline"
```

main.cpp:

```
#include "header.hpp"

int main() {}
```

source.cpp:

```
#include "header.hpp"

// some code
```

header.hpp:

```
template<typename T>
void f();
```

main.cpp:

```
#include "header.hpp"

int main() {
// f<int>(); // linking error
              // f<int> undefined
    f<char>(); // ok
}
```

source.cpp:

```
#include "header.hpp"

template<typename T>
void f() {} // valid in this
              // translation unit

void g() {
    f<char>(); // generates f<char>
}
```

Class Template

header.hpp:

```
template<typename T>
struct A {
    T    x = 3;
    void f() {}; // "inline" definition
};
```

main.cpp:

```
#include "header.hpp"

int main() {
    A<int>    a1; // ok
    A<float>  a2; // ok
    A<char>   a3; // ok
}
```

source.cpp:

```
#include "header.hpp"

int g() {
    A<int>    a1; // ok
    A<float>  a2; // ok
    A<char>   a3; // ok
}
```

header.hpp:

```
template<typename T>
struct A {
    T    x;
    void f(); // declaration
};
#include "header.i.hpp"
```

header.i.hpp:

```
template<typename T>
T A<T>::x = 3; // initialization

template<typename T>
void A<T>::f() {} // definition
```

main.cpp:

```
#include "header.hpp"

int main() {
    A<int>    a1; // ok
    A<float>  a2; // ok
    A<char>   a3; // ok
}
```

source.cpp:

```
#include "header.hpp"

int g() {
    A<int>    a1; // ok
    A<float>  a2; // ok
    A<char>   a3; // ok
}
```

Class Template Specialization - Case 1

header.hpp:

```
template<typename T>
struct A {
    T    x;
    void f(); // declaration
};
```

main.cpp:

```
#include "header.hpp"

int main() {
    A<int>  a1; // ok
    // A<char> a2; // linking error
}
```

source.cpp:

```
#include "header.hpp"

template<typename T>
int A<T>::x = 3; // initialization

template<typename T>
void A<T>::f() {} // definition

// template specialization
template class A<int>;
```

Class Template Specialization - Case 2

C++11

header.hpp:

```
template<typename T>
struct A {
    T    x;
    void f();
};
```

source.cpp:

```
#include "header.hpp"

extern template class A<int>;
// A<int> is not generated by the
// compiler in this translation unit
int main() {
    A<int> a;
}
```

source.cpp:

```
#include "header.hpp"

// template specialization
template class A<int>;

// or any instantiation of
// A<int>, e.g.  A<int> a;
```


Undefined Behavior and Summary

Undefined Behavior - inline Function

main.cpp:

```
#include <iostream>
inline int f() { return 3; }

void g();

int main() {
    std::cout << f(); // print 3
    std::cout << g(); // print 3!!
}                      // not 5
```

source.cpp:

```
// same signature and inline
inline int f() { return 5; }

int g() { return f(); }
```

The linker can *arbitrary* choose one of the two definitions of `f()` . With `-O3` , the compiler could *inline* `f()` in `g()` , so now `g()` return `5`

This issue is easy to detect in trivial examples but hard to find in large codebase

Solution: `static` or `anonymous namespace`

Undefined Behavior - Member Function

main.cpp:

```
#include <iostream>

struct A {
    int f() { return 3; }
};

int g();
```

main.cpp:

```
#include "header.hpp"
using namespace std;

int main() {
    A a;
    cout << a.f(); // print 3
    cout << g();  // print 3!!
}
```

source.cpp:

```
struct A {
    int f() { return 5; }
};

int g() {
    A<int> a;
    return a.f();
}
```

Undefined Behavior - Function Template

main.cpp:

```
template<typename T>
int f() {
    return 3;
}

int g();
```

source1.cpp:

```
#include "header.hpp"
using namespace std;

int main() {
    cout << f<int>(); // print 3
    cout << g();      // print 3!!
}
```

source2.cpp:

```
template<typename T>
int f() {
    return 5;
}

int g() {
    return f<int>();
}
```

Undefined Behavior

Other ODR violations are even harder (if not impossible) to find, e.g. Diagnosing Hidden ODR Violations in Visual C++

Some tools for partially detecting ODR violations:

- `-detect-odr-violations` flag for gold/llvm linker
- `-Wodr -flto` flag for GCC
- Clang address sanitizer + `ASAN_OPTIONS=detect_odr_violation=2`
(link)

Another solution could be include all files in a single translation unit

Summary

- **header:** declaration of
 - functions, structures, classes, types, alias
 - `template` function, structs, classes
 - `extern` variables, functions
 - global `const/constexpr` variables
- **header implementation:** definition of
 - `inline` functions, variables
 - `template` functions, classes
- **source file:** definition of
 - functions
 - `template` full specialization functions, classes
 - `static` global variables (+ declaration)
 - `extern` variables, functions

#include Issues

Forward Declaration

Forward declaration is a declaration of an identifier for which a complete definition has not yet given

“*forward*” means that an entity is declared before it is used

main.cpp:

```
void f(); // function forward declaration

class A; // class forward declaration

int main() {
    f(); // ok, f() is a function and not a variable
    // A a; // compiler error no definition (incomplete type)
           // e.g. the compiler is not able to deduce the size of A
    A* a; // ok
}
```

source.cpp:

```
void f() {} // definition of f()
class A {}; // definition of A()
```


Forward Declaration vs. `#include`

Advantages:

- Forward declarations can save compile time as `#include` forces the compiler to open more files and process more input
- Forward declarations can save on unnecessary recompilation. `#include` can force your code to be recompiled more often, due to unrelated changes in the header

Disadvantages:

- Forward declarations can hide a dependency, allowing user code to skip necessary recompilation when headers change
- A forward declaration may be broken by subsequent changes to the library
- Forward declaring multiple symbols from a header can be more verbose than simply `#including` the header

The `include guard` avoids the problem of multiple inclusions of a header file in a translation unit

`header.hpp`:

```
#ifndef HEADER_HPP // include guard
#define HEADER_HPP

... many lines of code ...


#endif // HEADER_HPP
```

`#pragma once` preprocessor directive is an alternative to the `include guard` to force current file to be included only once in a translation unit

- `#pragma once` is less portable but less verbose and compile faster than the `include guard`

The `include guard`/`#pragma once` should be used in every header file

Common case:



header_A.hpp:

```
#pragma once // prevents "multiple definitions" linking error
```

```
struct A {  
};
```

header_B.hpp:

```
#include "header_A.hpp" // included here
```

```
struct B {  
    A a;  
};
```


main.cpp:

```
#include "header_A.hpp" // .. and included here
```

```
#include "header_B.hpp"
```

```
int main() {  
    A a; // ok, here we need "header_A.hpp"  
    B b; // ok, here we need "header_B.hpp"  
}
```

A **circular dependency** is a relation between two or more modules which either directly or indirectly depend on each other to function properly



Circular dependencies can be solved by using forward declaration, or better, by rethinking the project organization

header_A.hpp:

```
#pragma once
#include "header_B.hpp"
class A {
    B* b;
};
```

header_B.hpp:

```
#pragma once
#include "header_C.hpp"
class B {
    C* c;
};
```

header_C.hpp:

```
#pragma once
#include "header_A.hpp"
class C { // compile error "header_A" already included by "main.cpp"
    A* a; // the compiler cannot view the "class C"
};
```

header_A.hpp:

```
#pragma once
class B;    // forward declaration
           // note: does not include "header_B.hpp"

class A {
    B* b;
};
```

header_B.hpp:

```
#pragma once
class C;    // forward declaration
class B {
    C* c;
};
```

header_C.hpp:

```
#pragma once
class A;    // forward declaration
class C {
    A* a;
};
```

Common Linking Errors

Very common *linking* errors:

- **undefined reference**

Solutions:

- Check if the right headers and sources are included
- Break circular dependencies

- **multiple definitions**

Solutions:

- **inline** function, variable definition or **extern** declaration
- Add include guard/ **#pragma once** to header files
- Place template definition in header file and full specialization in source files

Namespace

The problem: Named entities, such as variables, functions, and compound types declared outside any block has *global scope*, meaning that its name is valid anywhere in the code

Namespaces allow to group named entities that otherwise would have global scope into narrower scopes, giving them ***namespace scope*** (where *std* stands for “standard”)

Namespaces provide a method for preventing name conflicts in large projects. Symbols declared inside a namespace block are placed in a named scope that prevents them from being mistaken for identically-named symbols in other scopes

Namespace Functions vs. Class + static Methods

Namespace functions:

- Namespace can be extended anywhere (without control)
- Namespace specifier can be avoided with the keyword `using`

Class + `static` methods:

- Can interact only with static data members
- `struct/class` cannot be extended outside their declarations

`static` methods should define operations strictly related to object definition, otherwise `namespace` should be preferred

Defining a Namespace

```
#include <iostream>

namespace ns1 {
void f() {
    std::cout << "ns1" << std::endl;
}
} // namespace ns1

namespace ns2 {
void f() {
    std::cout << "ns2" << std::endl;
}
} // namespace ns2

int main () {
    ns1::f(); // print "ns1"
    ns2::f(); // print "ns2"
    // f();    // compile error f() is not visible
}
```

Namespace Conflicts

```
#include <iostream>
using namespace std;
void f() {
    cout << "global" << endl;
}
namespace ns1 {
    void f() { cout << "ns1::f()" << endl; }
    void g() { cout << "ns1::g()" << endl; }
}

int main () {
    f();          // ok, print "global"
    // g();       // compile error g() is not visible

    using namespace ns1;
    // f();       // compile error ambiguous function name
    ::f();        // ok, print "global"
    ns1::f();     // ok, print "ns1::f()"
    g();          // ok, print "ns1::g()", only one choice
}
```

Nested Namespaces

```
#include <iostream>
using namespace std;
namespace ns1 {
    void f() { cout << "ns1::f()" << endl; }
namespace ns2 {
    void f() { cout << "ns1::ns2::f()" << endl; }
}
}

namespace ns1 { // the same namespace can be declared multiple times,
namespace ns2 { // and extended in multiple files
    void g() {}
}
}
```

C++17 allows nested namespace definitions:

```
namespace ns1::ns2 {
    void h()
}
```

Namespace Scope

```
#include <iostream>
using namespace std;
namespace ns1 {
    void f() { cout << "ns1::f()" << endl; }
}
namespace ns2 {
    void f() { cout << "ns1::ns2::f()" << endl; }
    void g() { cout << "ns1::ns2::g()" << endl; }
}
namespace ns1 {
    void g() {} // ok
    // void f() {} // compile error function name conflict with
    //                                     header.hpp: "ns1::f()"
}

int main() {
    ns1::f(); // ok, print "ns1::f()"
    ns1::ns2::f(); // ok, print "ns1::ns2::f()"
    using namespace ns1::ns2;
    g(); // ok, print "ns1::ns2::g()"
}
```

Namespace Alias

Namespace alias allows declaring an alternate name for an existing namespace

```
namespace very_very_long_namespace {  
    void g() {}  
}  
  
int main() {  
    namespace ns = very_very_long_namespace; // namespace alias  
    ns::g();  
}
```


Anonymous Namespace

A namespace with no identifier before an opening brace produces an **unnamed/anonymous namespace**

Entities inside an anonymous namespace are used for declaring unique identifiers, visible in the same source file

Anonymous namespaces vs. static global entities

- Anonymous namespaces allow *type declarations*, and they are *less verbose*

main.cpp

```
#include <iostream>
namespace { // anonymous
    void f() { std::cout << "main"; }
}           // internal linkage

int main() {
    f();    // ok, print "main"
}
```

source.cpp

```
#include <iostream>
namespace { // anonymous
    void f() { std::cout << "source"; }
}

int g() {
    f();    // ok, print "source"
}
```

inline Namespace

inline namespace is a concept similar to library versioning. It is a mechanism that makes a nested namespace look and act as if all its declarations were in the surrounding namespace

```
namespace ns1 {  
    inline namespace V99 {  
        void f(int) {}    // most recent version  
    }  
    namespace V98 {  
        void f(int) {}  
    }  
}  
  
using namespace ns1;  
  
int main() {  
    V98::f(1);    // call V98  
    V99::f(1);    // call V99  
    f(1);         // call default version (V99)  
}
```

How to Compile

Compile Methods

Method 1

Compile all files together (naive):

```
g++ -I include/ main.cpp source.cpp -o main.x
```

-I: Specify the **include path** to the compiler. It can be used multiple times

Method 2

Compile each *translation unit* in a file object:

```
g++ -c -I include/ source.cpp -o source.o
```

```
g++ -c -I include/ main.cpp -o main.o
```

Link all file objects:

```
g++ main.o source.o -o main.x
```

A **library** is a package of code that is meant to be reused by many programs

A **static library** (.a) is a set of object files that are directly linked into the final executable. If a program is compiled with a static library, all the functionality of the static library becomes part of final executable

- A static library cannot be modified without re-link the final executable
- Increase the size of the final executable
- The linker can optimize the final executable (*link time optimization*)

A **dynamic library**, also called a **shared library** (.so), consists of routines that are loaded into the application at run-time. If a program is compiled with a dynamic library, the library does not become part of final executable. It remains as a separate unit

- A dynamic library can be modified without re-link
- Dynamic library functions are called outside the executable
- Neither the linker, nor the compiler can optimize the code between shared libraries and the final executable

Compile with Libraries

Specify the **library path** (path where search for static/dynamic libraries) to the compiler: `g++ -L<library_path> main.cpp -o main`

-L can be used multiple times

Specify the **library name** (e.g. liblibrary.a) to the compiler:

```
g++ -llibrary main.cpp -o main
```

Linux/Unix Environmental variables:

- **LIBRARY_PATH** Specify the directories where search for *static* libraries at *compile-time*. Used by the compiler
- **LD_LIBRARY_PATH** Specify the directories where search for *dynamic/shared* libraries at *run-time*. Used by the program

Build Static/Dynamic Libraries

Static Library Creation

- Create object files for each translation unit (.cpp)
- Create the static library by using the **archiver** (**ar**) linux utility

```
g++ source1.c -c source1.o
g++ source2.c -c source2.o
ar rvs libmystaticlib.a source1.o source2.o
```

Dynamic Library Creation

- Create object files for each translation unit (.cpp). Since library cannot store code at fixed addresses the compile must generate *position independent code*
- Create the dynamic library

```
g++ source1.c -c source1.o -fPIC
g++ source2.c -c source2.o -fPIC
g++ source1.o source2.o -shared -o libmydynamiclib.so
```


Find Dynamic Library Dependencies

The `ldd` utility shows the shared objects (shared libraries) required by a program or other shared objects

```
$ ldd /bin/ls
linux-vdso.so.1 (0x00007ffcc3563000)
libselinux.so.1 => /lib64/libselinux.so.1 (0x00007f87e5459000)
libcap.so.2 => /lib64/libcap.so.2 (0x00007f87e5254000)
libc.so.6 => /lib64/libc.so.6 (0x00007f87e4e92000)
libpcre.so.1 => /lib64/libpcre.so.1 (0x00007f87e4c22000)
libdl.so.2 => /lib64/libdl.so.2 (0x00007f87e4a1e000)
/lib64/ld-linux-x86-64.so.2 (0x00005574bf12e000)
libattr.so.1 => /lib64/libattr.so.1 (0x00007f87e4817000)
libpthread.so.0 => /lib64/libpthread.so.0 (0x00007f87e45fa000)
```

The `nm` utility provides information on the symbols being used in an object file or executable file

```
$ nm -D -C something.so
```

```
  w __gmon_start__
```

```
  D __libc_start_main
```

```
  D free
```

```
  D malloc
```

```
  D printf
```

```
# -C: Decode low-level symbol names
```

```
# -D: accepts a dynamic library
```

`readelf` displays information about ELF format object files

```
$ readelf --symbols something.so | c++filt
... OBJECT LOCAL DEFAULT 17 __frame_dummy_init_array_
... FILE LOCAL DEFAULT ABS prog.cpp
... OBJECT LOCAL DEFAULT 14 CC1
... OBJECT LOCAL DEFAULT 14 CC2
... FUNC LOCAL DEFAULT 12 g()
```

--symbols: display symbol table

`objdump` displays information about object files

```
$ objdump -t -C something.so | c++filt
... df *ABS*      ...  prog.cpp
...  0 .rodata    ...  CC1
...  0 .rodata    ...  CC2
...  F .text      ...  g()
...  0 .rodata    ...  (anonymous namespace)::CC3
...  0 .rodata    ...  (anonymous namespace)::CC4
...  F .text      ...  (anonymous namespace)::h()
...  F .text      ...  (anonymous namespace)::B::j1()
...  F .text      ...  (anonymous namespace)::B::j2()
```

--t: display symbols

-C: Decode low-level symbol names

References and Additional Material

- 20 ABI (Application Binary Interface) breaking changes every C++ developer should know
- Policies/Binary Compatibility Issues With C++
- 10 differences between static and dynamic libraries every C++ developer should know